

A TUTORIAL ON CONDITIONAL RANDOM FIELDS

Slim Essid

Telecom ParisTech

February 2015



Contents

- ▶ Introduction
- ▶ The logistic regression model
- ▶ Conditional Random Fields (for sequential data)
- ▶ Improvements and extensions to original CRFs
- ▶ Conclusion
- ▶ References
- ▶ Appendix

The supervised classification problem

Goal: predict *labels* y (aka *classes* or *outputs*) for some *observations* o (aka *data points*, *inputs*).

Examples:

- Predict *translation/part-of-speech tag* for a word.
- Predict *instrument, chord, notes* played... for a music segment.

Supervised classification:

- Each observation o is supposed to pertain to a **predefined** class \mathcal{C}_k : the k -th (**discrete**) class of a classification problem; $k = 1, \dots, K$.
- This is represented using a label y for each o ; $y \in \mathcal{Y}$, e.g. $\mathcal{Y} = \{0, 1\}$, $\mathcal{Y} = \{1, 2, 3, \dots, K\}$.

Examples of classes:

- POS tags: *noun, verb, adjective*...
- Music chords: *C7, Gmaj7, Fmin7*...

The supervised classification problem

Goal: predict *labels* y (aka *classes* or *outputs*) for some *observations* o (aka *data points*, *inputs*).

Examples:

- Predict *translation/part-of-speech tag* for a word.
- Predict *instrument, chord, notes* played... for a music segment.

Supervised classification:

- Each observation o is supposed to pertain to a **predefined** class \mathcal{C}_k : the k -th (**discrete**) class of a classification problem; $k = 1, \dots, K$.
- This is represented using a label y for each o ; $y \in \mathcal{Y}$, e.g. $\mathcal{Y} = \{0, 1\}$, $\mathcal{Y} = \{1, 2, 3, \dots, K\}$.


Examples of classes:

- **POS tags:** *noun, verb, adjective*...
- **Music chords:** *C7, Gmaj7, Fmin7*...

Features

- Classification relies on **features** x : descriptors of some qualities/attributes of the inputs o . Two types of features:

Continuous features



real-valued: e.g. MFCC, chroma, tempo...

Discrete/categorical




symbolic: e.g. note/chord played, key...

- Usually assembled as feature vectors x .

Features

- Classification relies on **features** x : descriptors of some qualities/attributes of the inputs o . Two types of features:

Continuous features



real-valued: e.g. MFCC, chroma, tempo...

Discrete/categorical



symbolic: e.g. note/chord played, key...

- Usually assembled as **feature vectors** x .

Notations

- o : an input (observation) to be classified; *e.g.: a word, an image, an audio frame/segment...*
- $\mathbf{x} = (x_1, \dots, x_D)^T$: a D -dimensional column vector (usually in \mathbb{R}^D); \mathbf{x}^T is a row vector.
- \mathbf{x}_n is a **feature vector** among a collection of N examples $\mathbf{x}_1, \dots, \mathbf{x}_N$.
- x_{jn} is the j -th **feature coefficient** of \mathbf{x}_n ; $1 \leq j \leq D$.
- $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$: the set of all training feature-vector examples.

Feature functions

Different from features!

Definition

A **feature function** is a real-valued function of both the input space \mathcal{O} (observations) and the output space \mathcal{Y} (target labels), $f_j : \mathcal{O} \times \mathcal{Y} \rightarrow \mathbb{R}$, that can be used to compute characteristics of the observations.

- An alternative way to express the characteristics of the observations, in a more **flexible manner**:
 - using output-specific features;
 - describing the **context**.

$$\text{Example: } f_j(o_i, y_i) = \begin{cases} 1 & \text{if } o_i = \text{"We"}, \text{ and } y_i = \text{"noun"} \\ 0 & \text{otherwise} \end{cases}$$

Feature functions

► Remarks:

- Different attributes may thus be considered for different classes.
- Feature functions are more general than features: one can define
 - $f_j(o, y) \triangleq x_j$;
 - or
 - $f_j(o, y) \triangleq x$.
- In the following:
 - Feature-function notations will be used only when needed.
 - Otherwise, feature-vectors will be preferred.

Feature functions

► Remarks:

- Different attributes may thus be considered for different classes.
- Feature functions are more general than features: one can define
 - $f_j(o, y) \triangleq x_j$;
 - or
 - $f_j(o, y) \triangleq \mathbf{x}$.
- In the following:
 - Feature-function notations will be used only when needed.
 - Otherwise, feature-vectors will be preferred.

Feature functions

► Remarks:

- Different attributes may thus be considered for different classes.
- Feature functions are more general than features: one can define
 - $f_j(o, y) \triangleq x_j$;
or
 - $f_j(o, y) \triangleq \mathbf{x}$.
- In the following:
 - Feature-function notations will be used only when needed.
 - Otherwise, feature-vectors will be preferred.

Probabilistic classification

Take decisions based on the **MAP** rule:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} p(y|\mathbf{x})$$

in order to minimize the error rate (here the expected 0-1 loss).

MAP: *Maximum A Posteriori* probability

→ this is the **Bayes decision rule** (for the 0-1 loss.)

How to get there?

Probabilistic classification

Take decisions based on the **MAP** rule:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} p(y|\mathbf{x})$$

in order to minimize the error rate (here the expected 0-1 loss).

MAP: *Maximum A Posteriori* probability

→ this is the **Bayes decision rule** (for the 0-1 loss.)

How to get there?

Generative model based classification

- **Objective:** $\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x})$.
- By the **Bayes** rule $p(y|\mathbf{x}) = \frac{p(y,\mathbf{x})}{p(\mathbf{x})} = \frac{p(y)p(\mathbf{x}|y)}{p(\mathbf{x})}$,

$$\hat{y} = \operatorname{argmax}_y \frac{p(y)p(\mathbf{x}|y)}{p(\mathbf{x})} = \operatorname{argmax}_y p(y)p(\mathbf{x}|y).$$

- Assuming a fixed prior $p(y)$ (possibly uninformative: $p(y) = \frac{1}{K}$), one is left with:

$$\hat{y} = \operatorname{argmax}_y p(\mathbf{x}|y).$$

- Our decision criterion becomes a **maximum-likelihood** criterion.
- This is a **generative approach** to classification: a probabilistic model of “how to generate \mathbf{x} given a class y ” is targeted.

Generative model based classification

- **Objective:** $\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x})$.
- By the **Bayes** rule $p(y|\mathbf{x}) = \frac{p(y,\mathbf{x})}{p(\mathbf{x})} = \frac{p(y)p(\mathbf{x}|y)}{p(\mathbf{x})}$,

$$\hat{y} = \operatorname{argmax}_y \frac{p(y)p(\mathbf{x}|y)}{p(\mathbf{x})} = \operatorname{argmax}_y p(y)p(\mathbf{x}|y).$$

- Assuming a fixed prior $p(y)$ (possibly uninformative: $p(y) = \frac{1}{K}$), one is left with:

$$\hat{y} = \operatorname{argmax}_y p(\mathbf{x}|y).$$

- Our decision criterion becomes a **maximum-likelihood** criterion.
- This is a **generative approach** to classification: a probabilistic model of “how to generate \mathbf{x} given a class y ” is targeted.

Discriminative model based classification



Directly models $p(y|\mathbf{x})$ without wasting efforts on modeling the observations, which is not needed for the goal $\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x})$.

Discriminative model based classification

Directly models $p(y|x)$ without wasting efforts on modeling the observations, which is not needed for the goal $\hat{y} = \operatorname{argmax}_y p(y|x)$.

► Pros:

- The class posteriors $p(y = c|x)$ are potentially simpler than the class-conditional densities.



Generated using pmtk3 (Dunham and Murphy, 2010)

Discriminative model based classification

Directly models $p(y|\mathbf{x})$ without wasting efforts on modeling the observations, which is not needed for the goal $\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x})$.

► Pros:

- The class posteriors $p(y = k|\mathbf{x})$ are potentially simpler than the class-conditional densities.
- Avoids making unwarranted assumptions about the features which may be highly **dependent** (especially with structured data).
- Improved robustness to model imperfections, as independence assumptions will be made only among the labels, not the observations.

► Cons:

- Classes need to be learned jointly and data should be available for all classes.
- Models do not allow for generating observations.

Discriminative model based classification

Directly models $p(y|\mathbf{x})$ without wasting efforts on modeling the observations, which is not needed for the goal $\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x})$.

► Pros:

- The class posteriors $p(y = k|\mathbf{x})$ are potentially simpler than the class-conditional densities.
- Avoids making unwarranted assumptions about the features which may be highly **dependent** (especially with structured data).
- Improved robustness to model imperfections, as independence assumptions will be made only among the labels, not the observations.

► Cons:

- Classes need to be learned jointly and data should be available for all classes.
- Models do not allow for generating observations.

Predicting structured-output data

- In many classification tasks the outputs are **structured**, e.g.:

Part-of-speech (POS) tagging:

tags follow predefined patterns


John	saw	the	big	table
<i>noun</i>	<i>verb</i>	<i>det</i>	<i>adj</i>	<i>noun</i>

Linear-chain structure

Predicting structured-output data

- In many classification tasks the outputs are **structured**, e.g.:

Musical instrument recognition




Linear-chain structure

Predicting structured-output data

- In many classification tasks the outputs are **structured**, e.g.:

Autotagging tasks:

target tags are correlated (e.g. comedy, romance, humour)



Tree structure


→ Need for predictors able to take advantage of this structure.

Predicting structured-output data

- In many classification tasks the outputs are **structured**, e.g.:

Autotagging tasks:

target tags are correlated (e.g. comedy, romance, humour)




Tree structure

→ Need for predictors able to take advantage of this structure.

Predicting sequential data

- In this tutorial, we focus on **sequential data**

John saw the big table
noun verb det adj noun



- Specialized inference algorithms can then be used (**forward-backward method**), which are easier to apprehend.
- More general methods can be used for more general structure (**belief propagation** and extensions), see for e.g. (Jensen and Nielsen, 2007).

More notations

- $\underline{\mathbf{x}}$ is a sequence of observations: $\underline{\mathbf{x}} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$.
- \underline{y} is the corresponding sequence of labels: $\underline{y} = (y_1, \dots, y_n)$.
- We assume we have a training dataset \mathcal{D} of N (i.i.d) such sequences:
 $\mathcal{D} = \{(\underline{\mathbf{x}}^{(1)}, \underline{y}^{(1)}), \dots, (\underline{\mathbf{x}}^{(N)}, \underline{y}^{(N)})\}$.
- Remarks:
 - Observations are no longer assumed to be i.i.d within each sequence.
 - Sequences $\underline{\mathbf{x}}^{(q)}$ do not necessarily have the same length, when needed n_q will denote the length of $\underline{\mathbf{x}}^{(q)}$.

More notations

- $\underline{\mathbf{x}}$ is a sequence of observations: $\underline{\mathbf{x}} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$.
- \underline{y} is the corresponding sequence of labels: $\underline{y} = (y_1, \dots, y_n)$.
- We assume we have a training dataset \mathcal{D} of N (i.i.d) such sequences:
 $\mathcal{D} = \{(\underline{\mathbf{x}}^{(1)}, \underline{y}^{(1)}), \dots, (\underline{\mathbf{x}}^{(N)}, \underline{y}^{(N)})\}$.
- **Remarks:**
 - Observations are no longer assumed to be i.i.d within each sequence.
 - Sequences $\underline{\mathbf{x}}^{(q)}$ do not necessarily have the same length, when needed n_q will denote the length of $\underline{\mathbf{x}}^{(q)}$.

The CRF model

A discriminative model for structured-output data

CRF model definition

$$\begin{aligned} p(\underline{y}|\underline{x}; \boldsymbol{\theta}) &= \frac{1}{Z(\underline{x}, \boldsymbol{\theta})} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y}) \\ &= \frac{1}{Z(\underline{x}, \boldsymbol{\theta})} \exp \Psi(\underline{x}, \underline{y}; \boldsymbol{\theta}); \quad \boldsymbol{\theta} = \{\theta_1, \dots, \theta_D\}. \end{aligned}$$

- $Z(\underline{x}, \boldsymbol{\theta}) = \sum_{\underline{y}} \exp \sum_j \theta_j F_j(\underline{x}, \underline{y})$ is called a **partition function**.
- $\Psi(\underline{x}, \underline{y}; \boldsymbol{\theta}) = \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y})$ is called a **potential function**.
- **Remark:** feature functions $F_j(\underline{x}, \underline{y})$ depend on the whole sequence of observations \underline{x} and labels \underline{y} .

The CRF model

A discriminative model for structured-output data

CRF model definition

$$\begin{aligned} p(\underline{y}|\underline{x}; \boldsymbol{\theta}) &= \frac{1}{Z(\underline{x}, \boldsymbol{\theta})} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y}) \\ &= \frac{1}{Z(\underline{x}, \boldsymbol{\theta})} \exp \Psi(\underline{x}, \underline{y}; \boldsymbol{\theta}); \quad \boldsymbol{\theta} = \{\theta_1, \dots, \theta_D\}. \end{aligned}$$

- $Z(\underline{x}, \boldsymbol{\theta}) = \sum_{\underline{y}} \exp \sum_j \theta_j F_j(\underline{x}, \underline{y})$ is called a **partition function**.
- $\Psi(\underline{x}, \underline{y}; \boldsymbol{\theta}) = \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y})$ is called a **potential function**.
- **Remark:** feature functions $F_j(\underline{x}, \underline{y})$ depend on the whole sequence of observations \underline{x} and labels \underline{y} .

Applications of CRFs

CRF models have proven to be superior to competitors in a variety of application fields.

- They are the state-of-the-art techniques in many **natural language processing** (NLP) tasks (Taskar et al., 2002; Settles, 2004; Lavergne et al., 2010)
part-of-speech tagging (POS), named-entity recognition (NER)...
- They have been successfully used for various **computer vision** tasks (He et al., 2004; Quattoni et al., 2007; Wang et al., 2006; Morency et al., 2007; Rudovic et al., 2012)
image labeling, object and gesture recognition, facial expressions...
- Also for **speech analysis** tasks (Gunawardana et al., 2005; Reiter et al., 2007; Morris and Fosler-Lussier, 2008; Hong, 2010)
speech recognition, speech segmentation, speaker identification...
- And a few applications to **music analysis** (Li, 2011; Schmidt and Kim, 2011; Imbrasaite et al., 2013; Joder et al., 2011).
autotagging, musical mood recognition, audio-to-score alignment, beat detection...

► Introduction

► The logistic regression model

- Model specification
- Maximum Entropy Modeling
- Parameter estimation
- Improvements to the logistic regression model

► Conditional Random Fields (for sequential data)

► Improvements and extensions to original CRFs

► Conclusion

► References

The logistic regression model

Approach: model the **posterior** probabilities of the K classes using linear functions of the inputs \mathbf{x} , according to:

$$\begin{aligned}\log \frac{P(C_1|\mathbf{x})}{P(C_K|\mathbf{x})} &= w_{10} + \mathbf{w}_1^T \mathbf{x} \\ \log \frac{P(C_2|\mathbf{x})}{P(C_K|\mathbf{x})} &= w_{20} + \mathbf{w}_2^T \mathbf{x} \\ &\vdots \\ \log \frac{P(C_{K-1}|\mathbf{x})}{P(C_K|\mathbf{x})} &= w_{(K-1)0} + \mathbf{w}_{K-1}^T \mathbf{x}\end{aligned}$$

Defines a **log-linear** model specified in terms of $K - 1$ **log-odds** ($\log \frac{P(C_k|\mathbf{x})}{P(C_K|\mathbf{x})}$) or **logit** transformations so that the K probabilities sum to 1.

The logistic regression model

- From $\log \frac{P(C_k|\mathbf{x})}{P(C_K|\mathbf{x})} = w_{k0} + \mathbf{w}_k^T \mathbf{x}$; $k = 1, \dots, K - 1$; it is easy to deduce that:

Multiclass logistic regression model

$$P(C_k|\mathbf{x}) = \frac{\exp(w_{k0} + \mathbf{w}_k^T \mathbf{x})}{1 + \sum_{l=1}^{K-1} \exp(w_{l0} + \mathbf{w}_l^T \mathbf{x})} ; k = 1, \dots, K - 1,$$

$$P(C_K|\mathbf{x}) = \frac{1}{1 + \sum_{l=1}^{K-1} \exp(w_{l0} + \mathbf{w}_l^T \mathbf{x})}$$

- Remarks
 - The model is a **classification** model (not a regression model!)
 - It is a **discriminative** model as it targets $P(C_k|\mathbf{x})$ (as opposed to modeling $p(\mathbf{x}|C_k)$ in **generative** models.)

The logistic regression model

- From $\log \frac{P(C_k|\mathbf{x})}{P(C_K|\mathbf{x})} = w_{k0} + \mathbf{w}_k^T \mathbf{x}$; $k = 1, \dots, K - 1$; it is easy to deduce that:

Multiclass logistic regression model

$$P(C_k|\mathbf{x}) = \frac{\exp(w_{k0} + \mathbf{w}_k^T \mathbf{x})}{1 + \sum_{l=1}^{K-1} \exp(w_{l0} + \mathbf{w}_l^T \mathbf{x})} ; k = 1, \dots, K - 1,$$

$$P(C_K|\mathbf{x}) = \frac{1}{1 + \sum_{l=1}^{K-1} \exp(w_{l0} + \mathbf{w}_l^T \mathbf{x})}$$

- Remarks
 - The model is a **classification** model (not a regression model!)
 - It is a **discriminative** model as it targets $P(C_k|\mathbf{x})$ (as opposed to modeling $p(\mathbf{x}|C_k)$ in **generative** models.)

Binary classification case

- When $K = 2$

$$\begin{aligned}P(\mathcal{C}_1|\mathbf{x}) &= p = \frac{1}{1 + \exp -(w_{10} + \mathbf{w}_1^T \mathbf{x})} \\P(\mathcal{C}_2|\mathbf{x}) &= 1 - p\end{aligned}$$

- $p = \frac{1}{1 + \exp -a}$; $a = w_{10} + \mathbf{w}_1^T \mathbf{x}$

Logistic sigmoid function

$$\sigma(a) \triangleq \frac{1}{1 + \exp -a}$$


The logistic sigmoid function

$$\sigma(a) \triangleq \frac{1}{1 + \exp -a}$$

- Properties:

Symmetry: $\sigma(-a) = 1 - \sigma(a)$

Inverse: $a = \log \frac{\sigma}{1-\sigma}$: **logit** function



- The odds $\frac{p}{1-p} \in [0, +\infty]$ hence the **log-odds** $\log \frac{p}{1-p} \in [-\infty, +\infty]$
- Logistic regression models the **log-odds** as **linear functions** of the inputs... *why is this a good idea?*


The logistic sigmoid function

$$\sigma(a) \triangleq \frac{1}{1 + \exp -a}$$

- Properties:

Symmetry: $\sigma(-a) = 1 - \sigma(a)$

Inverse: $a = \log \frac{\sigma}{1-\sigma}$: **logit** function



- The odds $\frac{p}{1-p} \in [0, +\infty]$ hence the **log-odds** $\log \frac{p}{1-p} \in [-\infty, +\infty]$
- Logistic regression models the **log-odds** as linear functions of the inputs... *why is this a good idea?*


The logistic sigmoid function

$$\sigma(a) \triangleq \frac{1}{1 + \exp -a}$$

- Properties:

Symmetry: $\sigma(-a) = 1 - \sigma(a)$

Inverse: $a = \log \frac{\sigma}{1-\sigma}$: **logit** function



- The odds $\frac{p}{1-p} \in [0, +\infty]$ hence the **log-odds** $\log \frac{p}{1-p} \in [-\infty, +\infty]$
- Logistic regression models the **log-odds** as **linear functions** of the inputs... *why is this a good idea?*

Maximum Entropy: an introductory example

By (Berger et al., 1996)

Goal: perform English-to-French translation.

Approach: model an expert translator's approach to decide to translate a particular word, e.g. "*in*".

Method: Use a training dataset to estimate $p(y|o)$: the probability to assign the word (or phrase) y to the observed word "*in*"; to be used for MAP decision.

- The structure of both the ground-truth labels y and the observations o reflect a set of facts about the data: the translator's expertise.

→ Our model should capture these facts to perform accurate predictions.

Maximum Entropy: an introductory example

By (Berger et al., 1996)

Goal: perform English-to-French translation.

Approach: model an expert translator's approach to decide to translate a particular word, e.g. "*in*".

Method: Use a training dataset to estimate $p(y|o)$: the probability to assign the word (or phrase) y to the observed word "*in*"; to be used for MAP decision.

- The structure of both the ground-truth labels y and the observations o reflect a set of **facts** about the data: the translator's expertise.

→ Our model should capture these facts to perform accurate predictions.

Using facts about the data

- *The translator always chooses among $\{dans, en, \grave{a}, au\}$ cours de, pendant $\}$:*

In terms of statistics

$$P(dans) + P(en) + P(\grave{a}) + P(au\ cours\ de) + P(pendant) = 1$$

- How to choose $P(dans), \dots, P(pendant)$?
- Safe choice:

In terms of statistics

$$P(dans) = P(en) = P(\grave{a}) = P(au\ cours\ de) = P(pendant) = 1/5$$

Using facts about the data

- *The translator always chooses among $\{dans, en, \grave{a}, au\}$ cours de, pendant $\}$:*

In terms of statistics

$$P(dans) + P(en) + P(\grave{a}) + P(au\ cours\ de) + P(pendant) = 1$$

- How to choose $P(dans), \dots, P(pendant)$?
- Safe choice:

In terms of statistics

$$P(dans) = P(en) = P(\grave{a}) = P(au\ cours\ de) = P(pendant) = 1/5$$

Using facts about the data

- *The translator always chooses among $\{dans, en, \grave{a}, au\}$ cours de, pendant $\}$:*

In terms of statistics

$$P(dans) + P(en) + P(\grave{a}) + P(au\ cours\ de) + P(pendant) = 1$$

- How to choose $P(dans), \dots, P(pendant)$?
- Safe choice:

In terms of statistics

$$P(dans) = P(en) = P(\grave{a}) = P(au\ cours\ de) = P(pendant) = 1/5$$

Why “uniform”?

- Intuitively: the most uniform model according to our knowledge, the only **unbiased** assumption
- Ancient wisdom:
 - *Occam's razor* (William of Ockham, 1287-1347): *principle of parsimony*: “*Nunquam ponenda est pluralitas sine necessitate.*” [Plurality must never be posited without necessity.]
 - *Laplace*: “when one has no information to distinguish between the probability of two events, the best strategy is to consider them equally likely.” (Principle of Insufficient Reason)

More facts

- *The translator chooses either “dans” or “en” 30% of the time:*

$$P(\text{dans}) + P(\text{en}) = 3/10$$

$$P(\text{dans}) + P(\text{en}) + P(\grave{\text{a}}) + P(\text{au cours de}) + P(\text{pendant}) = 1$$

- Again many solutions... and a reasonable choice is:

$$P(\text{dans}) = P(\text{en}) = 3/20$$

$$P(\grave{\text{a}}) = P(\text{au cours de}) = P(\text{pendant}) = 7/30$$

Even more facts

- The translator chooses either “dans” or “à” 50% of the time:

In terms of statistics

$$P(\text{dans}) + P(\text{en}) = 3/10$$

$$P(\text{dans}) + P(\text{en}) + P(\text{à}) + P(\text{au cours de}) + P(\text{pendant}) = 1$$

$$P(\text{dans}) + P(\text{à}) = 1/2$$

→ Less intuitive...

- What does “uniform” mean?
- How to determine the “most uniform” model subject to the constraints at hand?

Using feature functions

- Need to express the **facts** about the observations in a **flexible way**, to make sure the model will match them:
 - make use of **statistics** of the observations: e.g. “*in*” translates to either “*dans*” or “*en*” with frequency 3/10.
 - allow for using the **context**: e.g. if “*in*” is followed by “*April*” then the translation is “*en*” with frequency 9/10.

→ define **feature functions** to capture these statistics and use them to impose **constraints** to the model.

Example:
$$f_j(o_i, y_i) = \begin{cases} 1 & \text{if } y_i = \text{"en"} \text{ and "April" follows "in"} \\ 0 & \text{otherwise} \end{cases}$$

Using feature functions

- Need to express the **facts** about the observations in a **flexible way**, to make sure the model will match them:
 - make use of **statistics** of the observations: e.g. “*in*” translates to either “*dans*” or “*en*” with frequency 3/10.
 - allow for using the **context**: e.g. if “*in*” is followed by “*April*” then the translation is “*en*” with frequency 9/10.

→ define **feature functions** to capture these statistics and use them to impose **constraints** to the model.

$$\text{Example: } f_j(o_i, y_i) = \begin{cases} 1 & \text{if } y_i = \text{"en"} \text{ and "April" follows "in"} \\ 0 & \text{otherwise} \end{cases}$$

Defining constraints through feature functions

- The training sample can be described in terms of its **empirical** probability distribution $\tilde{p}(o, y)$:

$$\tilde{p}(o, y) \triangleq \frac{1}{N} \times \text{number of times that } (o, y) \text{ occurs in the sample}$$

- $\tilde{\mathbb{E}}(f_j) \triangleq \sum_{o, y} \tilde{p}(o, y) f_j(o, y)$: expected value of f_j w.r.t $\tilde{p}(o, y)$.
- $\mathbb{E}(f_j) \triangleq \sum_{o, y} p(o) p(y|o) f_j(o, y)$: expected v. of f_j w.r.t the **model** $p(o, y)$.

Defining constraints through feature functions

- The training sample can be described in terms of its **empirical** probability distribution $\tilde{p}(o, y)$:

$$\tilde{p}(o, y) \triangleq \frac{1}{N} \times \text{number of times that } (o, y) \text{ occurs in the sample}$$

- $\tilde{\mathbb{E}}(f_j) \triangleq \sum_{o, y} \tilde{p}(o, y) f_j(o, y)$: expected value of f_j w.r.t $\tilde{p}(o, y)$.
- $\mathbb{E}(f_j) \triangleq \sum_{o, y} p(o) p(y|o) f_j(o, y)$: expected v. of f_j w.r.t the **model** $p(o, y)$.

Defining constraints through feature functions

- The observed statistics (facts) are captured by enforcing:

Constraint equation

$$\mathbb{E}(f_j) = \tilde{\mathbb{E}}(f_j), \text{ i.e.}$$

$$\sum_{o,y} p(o)p(y|o)f_j(o,y) = \sum_{o,y} \tilde{p}(o,y)f_j(o,y)$$

- **Remark:** the constraints cannot be inconsistent since they are extracted from the data.

Defining constraints through feature functions

- The observed statistics (facts) are captured by enforcing:

Constraint equation

$$\mathbb{E}(f_j) = \tilde{\mathbb{E}}(f_j), \text{ i.e.}$$

$$\sum_{o,y} p(o)p(y|o)f_j(o,y) = \sum_{o,y} \tilde{p}(o,y)f_j(o,y)$$

- Remark:** the constraints cannot be inconsistent since they are extracted from the data.

Maximum entropy principle

- Now how to implement the idea of **uniform modeling**?
- Among the set \mathcal{M} of probability distributions that satisfy the **constraints**, $\mathbb{E}(f_j) = \tilde{\mathbb{E}}(f_j)$, choose:

Maximum entropy criterion

$$p^*(y|o) = \operatorname{argmax}_{p(y|o) \in \mathcal{M}} H(y|o);$$

$$H(y|o) \triangleq - \sum_{o,y} p(o)p(y|o) \log p(y|o) : \text{ the } \mathbf{conditional\ entropy}$$

- Hint from information theory: the discrete distribution with maximum **entropy** is the **uniform** distribution.

Solving the problem

Primal: $p^*(y|o) = \operatorname{argmax}_{p(y|o) \in \mathcal{M}} H(y|o)$

Constraints: $\mathbb{E}(f_j) = \tilde{\mathbb{E}}(f_j)$ and $\sum_y p(y|o) = 1$

Lagrangian: $L(p, \lambda) \triangleq H(y|o) + \lambda_0 \left(\sum_y p(y|o) - 1 \right) + \sum_j \lambda_j \left(\mathbb{E}(f_j) - \tilde{\mathbb{E}}(f_j) \right)$

Equating the derivative of the Lagrangian with 0:

$$p_{\lambda}(y|o) = \frac{1}{Z_{\lambda}(o)} \exp \sum_j \lambda_j f_j(o, y);$$

$$Z_{\lambda}(x) = \sum_y \exp \left(\sum_j \lambda_j f_j(o, y) \right)$$

The solution is given by the dual optimal: $\lambda^* = \operatorname{argmax}_{\lambda} L(p, \lambda).$

Solving the problem

Primal: $p^*(y|o) = \operatorname{argmax}_{p(y|o) \in \mathcal{M}} H(y|o)$

Constraints: $\mathbb{E}(f_j) = \tilde{\mathbb{E}}(f_j)$ and $\sum_y p(y|o) = 1$

Lagrangian: $L(p, \lambda) \triangleq H(y|o) + \lambda_0 \left(\sum_y p(y|o) - 1 \right) + \sum_j \lambda_j \left(\mathbb{E}(f_j) - \tilde{\mathbb{E}}(f_j) \right)$

Equating the derivative of the Lagrangian with 0:

$$p_{\lambda}(y|o) = \frac{1}{Z_{\lambda}(o)} \exp \sum_j \lambda_j f_j(o, y);$$

$$Z_{\lambda}(x) = \sum_y \exp \left(\sum_j \lambda_j f_j(o, y) \right)$$

The solution is given by the dual optimal: $\lambda^* = \operatorname{argmax}_{\lambda} L(p, \lambda)$.

Solving the problem

Primal: $p^*(y|o) = \operatorname{argmax}_{p(y|o) \in \mathcal{M}} H(y|o)$

Constraints: $\mathbb{E}(f_j) = \tilde{\mathbb{E}}(f_j)$ and $\sum_y p(y|o) = 1$

Lagrangian: $L(p, \lambda) \triangleq H(y|o) + \lambda_0 \left(\sum_y p(y|o) - 1 \right) + \sum_j \lambda_j \left(\mathbb{E}(f_j) - \tilde{\mathbb{E}}(f_j) \right)$

Equating the derivative of the Lagrangian with 0:

$$p_{\lambda}(y|o) = \frac{1}{Z_{\lambda}(o)} \exp \sum_j \lambda_j f_j(o, y);$$

$$Z_{\lambda}(x) = \sum_y \exp \left(\sum_j \lambda_j f_j(o, y) \right)$$

The solution is given by the dual optimal: $\lambda^* = \operatorname{argmax}_{\lambda} L(p, \lambda)$.

Solving the problem

Primal: $p^*(y|o) = \operatorname{argmax}_{p(y|o) \in \mathcal{M}} H(y|o)$

Constraints: $\mathbb{E}(f_j) = \tilde{\mathbb{E}}(f_j)$ and $\sum_y p(y|o) = 1$

Lagrangian: $L(p, \lambda) \triangleq H(y|o) + \lambda_0 \left(\sum_y p(y|o) - 1 \right) + \sum_j \lambda_j \left(\mathbb{E}(f_j) - \tilde{\mathbb{E}}(f_j) \right)$

Equating the derivative of the Lagrangian with 0:

$$p_{\lambda}(y|o) = \frac{1}{Z_{\lambda}(o)} \exp \sum_j \lambda_j f_j(o, y);$$

$$Z_{\lambda}(x) = \sum_y \exp \left(\sum_j \lambda_j f_j(o, y) \right)$$

The solution is given by the dual optimal: $\lambda^* = \operatorname{argmax}_{\lambda} L(p, \lambda)$.

Compare to the LR model

► Maxent model:

$$p(y = k|o) = \frac{1}{Z_{\lambda}(o)} \exp \left(\sum_j \lambda_{jk} f_j(o, y) \right);$$

$$Z_{\lambda}(o) = \sum_y \exp \left(\sum_j \lambda_{jk} f_j(o, y) \right).$$

► Logistic regression model:

$$\begin{aligned} p(y = k|\mathbf{x}) &= \frac{\exp(w_{k0} + \mathbf{w}_k^T \mathbf{x})}{1 + \sum_{l=1}^{K-1} \exp(w_{l0} + \mathbf{w}_l^T \mathbf{x})} \\ &= \frac{\exp(w'_{k0} + \mathbf{w}'_k^T \mathbf{x})}{\sum_{l=1}^K \exp(w'_{l0} + \mathbf{w}'_l^T \mathbf{x})} \\ &= \frac{1}{Z_{\mathbf{w}}(\mathbf{x})} \exp(w'_{k0} + \mathbf{w}'_k^T \mathbf{x}). \end{aligned}$$

Compare to the LR model

► Maxent model:

$$p(y = k|o) = \frac{1}{Z_{\lambda}(o)} \exp \sum_j \lambda_{kj} f_j(o, y)$$

► Logistic regression model:

$$p(y = k|\mathbf{x}) = \frac{1}{Z_{\mathbf{w}}(\mathbf{x})} \exp(w'_{k0} + \mathbf{w}'_k{}^T \mathbf{x})$$

- Using:

- feature-function: $f_j(o, y) = x_j$; $f_0(o, y) = 1$ and $\mathbf{x} = (x_1, \dots, x_j, \dots, x_D)^T$;
- $w'_{k0} + \mathbf{w}'_k{}^T \mathbf{x} = \sum_{j=0}^D w'_{kj} f_j(o, y)$;

Compare to the LR model

► Maxent model:

$$p(y = k|o) = \frac{1}{Z_{\lambda}(o)} \exp \sum_j \lambda_{kj} f_j(o, y)$$

► Logistic regression model:

$$p(y = k|\mathbf{x}) = \frac{1}{Z_{\mathbf{w}}(\mathbf{x})} \exp(w'_{k0} + \mathbf{w}'_k{}^T \mathbf{x})$$

- Using:

- feature-function: $f_j(o, y) = x_j$; $f_0(o, y) = 1$ and $\mathbf{x} = (x_1, \dots, x_j, \dots, x_D)^T$;
- $w'_{k0} + \mathbf{w}'_k{}^T \mathbf{x} = \sum_{j=0}^D w'_{kj} f_j(o, y)$;

Compare to the CRF model

► Maxent model:

$$p(y = k|o) = \frac{1}{Z_{\lambda}(o)} \exp \sum_j \lambda_{kj} f_j(o, y)$$

► Logistic regression model:

$$p(y = k|o) = \frac{1}{Z_{\mathbf{w}}(o)} \exp \sum_j w'_{kj} f_j(o, y)$$

► CRF model:

$$p(\underline{y}|\underline{x}; \theta) = \frac{1}{Z_{\theta}(\underline{x})} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y})$$

Compare to the CRF model

► Maxent model:

$$p(y = k|o) = \frac{1}{Z_{\lambda}(o)} \exp \sum_j \lambda_{kj} f_j(o, y)$$

► Logistic regression model:

$$p(y = k|o) = \frac{1}{Z_{\mathbf{w}}(o)} \exp \sum_j w'_{kj} f_j(o, y)$$

► CRF model:

$$p(\underline{y}|\underline{\mathbf{x}}; \boldsymbol{\theta}) = \frac{1}{Z_{\boldsymbol{\theta}}(\underline{\mathbf{x}})} \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y})$$

Conclusion

The solution to the maximum entropy models has the same parametric form as logistic regression and CRF models.

- It is easily shown that the optimal solution is the **maximum-likelihood** solution in the parametric family $p_{\lambda}(y|\mathbf{x}) = \frac{1}{Z_{\lambda}(\mathbf{x})} \exp(\sum_j \lambda_j x_j)$.
- We've only considered discrete inputs, what about **continuous** inputs?
 - It is found that if the class-conditional densities $p(\mathbf{x}|y)$ are members of the **exponential family** of distributions, then the posterior probabilities are again given by **logistic sigmoids** of a linear function.
 - In particular, the model is optimal with **Gaussian densities** (with a shared covariance matrix).

The logistic regression model is quite well justified in a variety of situations.

Conclusion

The solution to the maximum entropy models has the same parametric form as logistic regression and CRF models.

- It is easily shown that the optimal solution is the **maximum-likelihood** solution in the parametric family $p_{\lambda}(y|\mathbf{x}) = \frac{1}{Z_{\lambda}(\mathbf{x})} \exp(\sum_j \lambda_j x_j)$.
- We've only considered discrete inputs, what about **continuous** inputs?
 - It is found that if the class-conditional densities $p(\mathbf{x}|y)$ are members of the **exponential family** of distributions, then the posterior probabilities are again given by **logistic sigmoids** of a linear function.
 - In particular, the model is optimal with **Gaussian densities** (with a shared covariance matrix).

The logistic regression model is quite well justified in a variety of situations.

Conclusion

The solution to the maximum entropy models has the same parametric form as logistic regression and CRF models.

- It is easily shown that the optimal solution is the **maximum-likelihood** solution in the parametric family $p_{\lambda}(y|\mathbf{x}) = \frac{1}{Z_{\lambda}(\mathbf{x})} \exp(\sum_j \lambda_j x_j)$.
- We've only considered discrete inputs, what about **continuous** inputs?
 - It is found that if the class-conditional densities $p(\mathbf{x}|y)$ are members of the **exponential family** of distributions, then the posterior probabilities are again given by **logistic sigmoids** of a linear function.
 - In particular, the model is optimal with **Gaussian densities** (with a shared covariance matrix).

The logistic regression model is quite well justified in a variety of situations.

► Introduction

► The logistic regression model

- Model specification
- Maximum Entropy Modeling
- Parameter estimation
- Improvements to the logistic regression model

► Conditional Random Fields (for sequential data)

► Improvements and extensions to original CRFs

► Conclusion

► References

Fitting the LR models

- Done by **maximum likelihood** estimation; in practice minimizing the **Negative Log-Likelihood (NLL)**.
- Let θ denote the set of all parameters:
 $\theta = \{w_{10}, \mathbf{w}_1, \dots, w_{(K-1)0}, \mathbf{w}_{K-1}\}.$
- The **log-likelihood** for the N (i.i.d) feature-vector observations is:

$$L(\mathcal{D}; \theta) \triangleq - \sum_{i=1}^N \log p(y_i | \mathbf{x}_i; \theta)$$

- To simplify, we focus on the **bi-class** case...

Fitting the LR models

- Done by **maximum likelihood** estimation; in practice minimizing the **Negative Log-Likelihood (NLL)**.
- Let θ denote the set of all parameters:
 $\theta = \{w_{10}, \mathbf{w}_1, \dots, w_{(K-1)0}, \mathbf{w}_{K-1}\}$.
- The **log-likelihood** for the N (i.i.d) feature-vector observations is:

$$L(\mathcal{D}; \theta) \triangleq - \sum_{i=1}^N \log p(y_i | \mathbf{x}_i; \theta)$$

- To simplify, we focus on the **bi-class** case...

NLL for bi-class LR

- Let $y_i = 1$ for \mathcal{C}_1 observations and $y_i = 0$ for \mathcal{C}_2 observations.
- Let $p(\mathbf{x}; \boldsymbol{\theta}) \triangleq p(y_i = 1 | \mathbf{x}_i; \boldsymbol{\theta})$; hence $p(y_i = 0 | \mathbf{x}_i; \boldsymbol{\theta}) = 1 - p(\mathbf{x}; \boldsymbol{\theta})$.
- We can write: $p(y | \mathbf{x}; \boldsymbol{\theta}) = p(\mathbf{x}; \boldsymbol{\theta})^y (1 - p(\mathbf{x}; \boldsymbol{\theta}))^{1-y}$.

Negative Log-Likelihood

$$\begin{aligned} L(\mathcal{D}; \boldsymbol{\theta}) = L(\tilde{\mathbf{w}}) &= - \sum_{i=1}^N \{ y_i \log p(\mathbf{x}_i; \tilde{\mathbf{w}}) + (1 - y_i) \log (1 - p(\mathbf{x}_i; \tilde{\mathbf{w}})) \} \\ &= - \sum_{i=1}^N \left\{ y_i \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i - \log \left(1 + \exp(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i) \right) \right\} \end{aligned}$$

where $\tilde{\mathbf{w}} = (w_0, \mathbf{w})$ and $\tilde{\mathbf{x}}_i = (1, \mathbf{x}_i)$ so that $\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i = w_0 + \mathbf{w}^T \mathbf{x}_i$.

NLL for bi-class LR

- Let $y_i = 1$ for \mathcal{C}_1 observations and $y_i = 0$ for \mathcal{C}_2 observations.
- Let $p(\mathbf{x}; \boldsymbol{\theta}) \triangleq p(y_i = 1 | \mathbf{x}_i; \boldsymbol{\theta})$; hence $p(y_i = 0 | \mathbf{x}_i; \boldsymbol{\theta}) = 1 - p(\mathbf{x}; \boldsymbol{\theta})$.
- We can write: $p(y | \mathbf{x}; \boldsymbol{\theta}) = p(\mathbf{x}; \boldsymbol{\theta})^y (1 - p(\mathbf{x}; \boldsymbol{\theta}))^{1-y}$.

Negative Log-Likelihood

$$\begin{aligned} L(\mathcal{D}; \boldsymbol{\theta}) = L(\tilde{\mathbf{w}}) &= - \sum_{i=1}^N \{y_i \log p(\mathbf{x}_i; \tilde{\mathbf{w}}) + (1 - y_i) \log (1 - p(\mathbf{x}_i; \tilde{\mathbf{w}}))\} \\ &= - \sum_{i=1}^N \left\{ y_i \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i - \log \left(1 + \exp(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i) \right) \right\} \end{aligned}$$

where $\tilde{\mathbf{w}} = (w_0, \mathbf{w})$ and $\tilde{\mathbf{x}}_i = (1, \mathbf{x}_i)$ so that $\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i = w_0 + \mathbf{w}^T \mathbf{x}_i$.

Gradient and Hessian of the NLL

Gradient: $\nabla L(\mathcal{D}; \tilde{\mathbf{w}}) = - \sum_{i=1}^N \tilde{\mathbf{x}}_i (y_i - p(\mathbf{x}_i; \tilde{\mathbf{w}}))$

Hessian: $\frac{\partial^2 L(\mathcal{D}; \tilde{\mathbf{w}})}{\partial \tilde{\mathbf{w}} \partial \tilde{\mathbf{w}}^T} = \sum_{i=1}^N \tilde{\mathbf{x}}_i \tilde{\mathbf{x}}_i^T p(\mathbf{x}_i; \tilde{\mathbf{w}}) (1 - p(\mathbf{x}_i; \tilde{\mathbf{w}}))$

- so the Hessian is **positive semi-definite**,
- the NLL is **convex** and it has a global minimum.

Gradient and Hessian of the NLL

Gradient: $\nabla L(\mathcal{D}; \tilde{\mathbf{w}}) = - \sum_{i=1}^N \tilde{\mathbf{x}}_i (y_i - p(\mathbf{x}_i; \tilde{\mathbf{w}}))$

Hessian: $\frac{\partial^2 L(\mathcal{D}; \tilde{\mathbf{w}})}{\partial \tilde{\mathbf{w}} \partial \tilde{\mathbf{w}}^T} = \sum_{i=1}^N \tilde{\mathbf{x}}_i \tilde{\mathbf{x}}_i^T p(\mathbf{x}_i; \tilde{\mathbf{w}}) (1 - p(\mathbf{x}_i; \tilde{\mathbf{w}}))$

- so the Hessian is **positive semi-definite**,
- the NLL is **convex** and it has a global minimum.

Gradient and Hessian of the NLL

Gradient: $\nabla L(\mathcal{D}; \tilde{\mathbf{w}}) = - \sum_{i=1}^N \tilde{\mathbf{x}}_i (y_i - p(\mathbf{x}_i; \tilde{\mathbf{w}}))$

Hessian: $\frac{\partial^2 L(\mathcal{D}; \tilde{\mathbf{w}})}{\partial \tilde{\mathbf{w}} \partial \tilde{\mathbf{w}}^T} = \sum_{i=1}^N \tilde{\mathbf{x}}_i \tilde{\mathbf{x}}_i^T p(\mathbf{x}_i; \tilde{\mathbf{w}}) (1 - p(\mathbf{x}_i; \tilde{\mathbf{w}}))$

- so the Hessian is **positive semi-definite**,
- the NLL is **convex** and it has a global minimum.

Minimizing the NLL

- By setting the derivatives to zero:

$$\frac{\partial L(\mathcal{D}; \tilde{\mathbf{w}})}{\partial w_j} = - \sum_{i=1}^N \tilde{x}_{ji} (y_i - p(\mathbf{x}_i; \tilde{\mathbf{w}})) = 0; 0 \leq j \leq D.$$

Optimization problem

Solve for $\tilde{\mathbf{w}}$ the $D + 1$ **non-linear** equations:

$$\sum_{i=1}^N y_i \tilde{x}_{ji} = \sum_{i=1}^N \tilde{x}_{ji} p(\mathbf{x}_i; \tilde{\mathbf{w}})$$

Optimization methods

Objective: Solve $\sum_{i=1}^N y_i \tilde{x}_{ji} = \sum_{i=1}^N \tilde{x}_{ji} p(\mathbf{x}_i; \tilde{\mathbf{w}})$

Problem: No closed-form solution in general (system of $D + 1$ non-linear equations).

Solution: use descent methods.

Among the many descent algorithms available, two are widely used:

- the **Newton-Raphson** method: fast... but complex (efficient variations exist);
- the **stochastic gradient** descent method (Bottou, 2010): easy to implement, adapted to large scale problems (*see Appendix*).

Optimization methods

Objective: Solve $\sum_{i=1}^N y_i \tilde{x}_{ji} = \sum_{i=1}^N \tilde{x}_{ji} p(\mathbf{x}_i; \tilde{\mathbf{w}})$

Problem: No closed-form solution in general (system of $D + 1$ **non-linear** equations).

Solution: use descent methods.

Among the many descent algorithms available, two are widely used:

- the **Newton-Raphson** method: fast... but complex (efficient variations exist);
- the **stochastic gradient** descent method (Bottou, 2010): easy to implement, adapted to large scale problems (*see Appendix*).

Optimization methods

Objective: Solve $\sum_{i=1}^N y_i \tilde{x}_{ji} = \sum_{i=1}^N \tilde{x}_{ji} p(\mathbf{x}_i; \tilde{\mathbf{w}})$

Problem: No closed-form solution in general (system of $D + 1$ **non-linear** equations).

Solution: use **descent** methods.

Among the many descent algorithms available, two are widely used:

- the **Newton-Raphson** method: fast... but complex (efficient variations exist);
- the **stochastic gradient** descent method (Bottou, 2010): easy to implement, adapted to large scale problems (see *Appendix*).

Optimization with the Newton-Raphson method

- To minimize $g(\theta)$, consider its second-order Taylor series approximation around θ_n :

$$g(\theta) \approx g(\theta_n) + \nabla g(\theta_n)^T (\theta - \theta_n) + \frac{1}{2} (\theta - \theta_n)^T H(\theta_n) (\theta - \theta_n);$$

$\nabla g(\theta_n)$ and $H(\theta_n)$ are resp. the **gradient** and **Hessian** of $g(\theta)$ at θ_n .

- This approximation is a quadratic function which is minimized by solving:

$$\nabla g(\theta_n) + H(\theta_n)(\theta - \theta_n) = 0.$$

Hence the **Newton-Raphson step**

$$\theta_{n+1} = \theta_n - H(\theta_n)^{-1} \nabla g(\theta_n).$$

Optimization with the Newton-Raphson method

- To minimize $g(\theta)$, consider its second-order Taylor series approximation around θ_n :

$$g(\theta) \approx g(\theta_n) + \nabla g(\theta_n)^T (\theta - \theta_n) + \frac{1}{2} (\theta - \theta_n)^T H(\theta_n) (\theta - \theta_n);$$

$\nabla g(\theta_n)$ and $H(\theta_n)$ are resp. the **gradient** and **Hessian** of $g(\theta)$ at θ_n .

- This approximation is a quadratic function which is minimized by solving:

$$\nabla g(\theta_n) + H(\theta_n)(\theta - \theta_n) = 0.$$

Hence the **Newton-Raphson step**

$$\theta_{n+1} = \theta_n - H(\theta_n)^{-1} \nabla g(\theta_n).$$

Optimization with the Newton-Raphson method

- To minimize $g(\theta)$, consider its second-order Taylor series approximation around θ_n :

$$g(\theta) \approx g(\theta_n) + \nabla g(\theta_n)^T (\theta - \theta_n) + \frac{1}{2} (\theta - \theta_n)^T H(\theta_n) (\theta - \theta_n);$$

$\nabla g(\theta_n)$ and $H(\theta_n)$ are resp. the **gradient** and **Hessian** of $g(\theta)$ at θ_n .

- This approximation is a quadratic function which is minimized by solving:


$$\nabla g(\theta_n) + H(\theta_n)(\theta - \theta_n) = 0.$$

Hence the **Newton-Raphson step**

$$\theta_{n+1} = \theta_n - H(\theta_n)^{-1} \nabla g(\theta_n).$$

Optimization with the Newton-Raphson method


Illustration



Generated using pmtk3 (Dunham and Murphy, 2010)

Optimization with the Newton-Raphson method

Illustration



Generated using pmtk3 (Dunham and Murphy, 2010)

Optimization with the Newton-Raphson method

Discussion

- Typically the algorithm converges (though overshooting may occur), and convergence speed is quadratic.
- D has to be small enough so that it is not too costly to **recompute** and **store** the inverse Hessian matrix at each iteration.
- Otherwise use **Quasi-Newton methods**:
 - **BFGS** (Broyden, Fletcher, Goldfarb and Shanno) method: approximates the inverse Hessian using successive gradient values.
 - **L-BFGS** (limited memory BFGS) method: stores only a few vectors used to approximate the inverse Hessian.
- Alternatively, use **stochastic gradient learning**:
 - Makes gradient updates based on **one** training example **at a time**.
 - In practice: simple approach, slow convergence, less accurate than L-BFGS.

Optimization with the Newton-Raphson method

Discussion

- Typically the algorithm converges (though overshooting may occur), and convergence speed is quadratic.
- D has to be small enough so that it is not too costly to **recompute** and **store** the inverse Hessian matrix at each iteration.
- Otherwise use **Quasi-Newton methods**:
 - **BFGS** (Broyden, Fletcher, Goldfarb and Shanno) method: approximates the inverse Hessian using successive gradient values.
 - **L-BFGS** (limited memory BFGS) method: stores only a few vectors used to approximate the inverse Hessian.
- Alternatively, use **stochastic gradient learning**:
 - Makes gradient updates based on **one** training example **at a time**.
 - In practice: simple approach, slow convergence, less accurate than L-BFGS.

Optimization with the Newton-Raphson method

Discussion

- Typically the algorithm converges (though overshooting may occur), and convergence speed is quadratic.
- D has to be small enough so that it is not too costly to **recompute** and **store** the inverse Hessian matrix at each iteration.
- Otherwise use **Quasi-Newton methods**:
 - **BFGS** (Broyden, Fletcher, Goldfarb and Shanno) method: approximates the inverse Hessian using successive gradient values.
 - **L-BFGS** (limited memory BFGS) method: stores only a few vectors used to approximate the inverse Hessian.
- Alternatively, use **stochastic gradient learning**:
 - Makes gradient updates based on **one** training example **at a time**.
 - In practice: simple approach, slow convergence, less accurate than L-BFGS.

► Introduction

► The logistic regression model

- Model specification
- Maximum Entropy Modeling
- Parameter estimation
- Improvements to the logistic regression model

► Conditional Random Fields (for sequential data)

► Improvements and extensions to original CRFs

► Conclusion

► References

ℓ_2 -regularization

- To **avoid overfitting** the complexity of the model should be penalized.
- Similarly to **ridge regression** (Hastie et al., 2009), a quadratic regularization term can be added to the NLL:

Regularized logistic regression problem

$$\begin{aligned}\hat{\mathbf{w}} &= \underset{\tilde{\mathbf{w}}}{\operatorname{argmin}} L(\mathcal{D}; \tilde{\mathbf{w}}) + \frac{\gamma}{2} \|\mathbf{w}\|^2 \\ &= \underset{\tilde{\mathbf{w}}}{\operatorname{argmin}} \left\{ - \sum_{i=1}^N \left[y_i \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i - \log \left(1 + \exp \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i \right) \right] + \frac{\gamma}{2} \sum_{j=1}^D w_j^2 \right\}\end{aligned}$$

$\gamma \geq 0$: complexity parameter controlling the amount of shrinkage; usually tuned by **cross-validation**.

ℓ_2 -regularization


Discussion

Recall that:

$$\hat{\mathbf{w}} = \underset{\tilde{\mathbf{w}}}{\operatorname{argmin}} L(\mathcal{D}; \tilde{\mathbf{w}}) + \frac{\gamma}{2} \|\mathbf{w}\|^2$$

is equivalent to:

$$\begin{cases} \hat{\mathbf{w}} = \underset{\tilde{\mathbf{w}}}{\operatorname{argmin}} L(\mathcal{D}; \tilde{\mathbf{w}}) \\ \text{subject to } \|\mathbf{w}\|^2 \leq t \end{cases}$$



for some t which has a correspondence to γ .

ℓ_2 -regularization

Gradient and Hessian

Gradient: $\nabla L_2(\mathcal{D}; \tilde{\mathbf{w}}) = \nabla L_2(\mathcal{D}; \tilde{\mathbf{w}}) + \gamma \mathbf{w}$

Hessian: $H_2(\tilde{\mathbf{w}}) = H(\tilde{\mathbf{w}}) + \gamma \mathbf{I}_{D+1}$

- So the Hessian becomes **positive definite**, the NLL is now **strictly convex** and it has a unique global minimum.
- The previous optimization methods can be straightforwardly adapted by modifying the expressions of the gradient and Hessian.

ℓ_2 -regularization

Gradient and Hessian

Gradient: $\nabla L_2(\mathcal{D}; \tilde{\mathbf{w}}) = \nabla L_2(\mathcal{D}; \tilde{\mathbf{w}}) + \gamma \mathbf{w}$

Hessian: $H_2(\tilde{\mathbf{w}}) = H(\tilde{\mathbf{w}}) + \gamma \mathbf{I}_{D+1}$

- So the Hessian becomes **positive definite**, the NLL is now **strictly convex** and it has a unique global minimum.
- The previous optimization methods can be straightforwardly adapted by modifying the expressions of the gradient and Hessian.

ℓ_2 -regularization

Gradient and Hessian

Gradient: $\nabla L_2(\mathcal{D}; \tilde{\mathbf{w}}) = \nabla L_2(\mathcal{D}; \tilde{\mathbf{w}}) + \gamma \mathbf{w}$

Hessian: $H_2(\tilde{\mathbf{w}}) = H(\tilde{\mathbf{w}}) + \gamma \mathbf{I}_{D+1}$

- So the Hessian becomes **positive definite**, the NLL is now **strictly convex** and it has a unique global minimum.
- The previous optimization methods can be straightforwardly adapted by modifying the expressions of the gradient and Hessian.

ℓ_2 -regularization

Gradient and Hessian

Gradient: $\nabla L_2(\mathcal{D}; \tilde{\mathbf{w}}) = \nabla L_2(\mathcal{D}; \tilde{\mathbf{w}}) + \gamma \mathbf{w}$

Hessian: $H_2(\tilde{\mathbf{w}}) = H(\tilde{\mathbf{w}}) + \gamma \mathbf{I}_{D+1}$

- So the Hessian becomes **positive definite**, the NLL is now **strictly convex** and it has a unique global minimum.
- The previous optimization methods can be straightforwardly adapted by modifying the expressions of the gradient and Hessian.

ℓ_1 -regularization

- Proceed as in the **LASSO** (Hastie et al., 2009), using a ℓ_1 -regularization.

ℓ_1 -regularized logistic regression problem

$$\begin{aligned}\hat{\mathbf{w}} &= \underset{\tilde{\mathbf{w}}}{\operatorname{argmin}} L(\mathcal{D}; \tilde{\mathbf{w}}) + \gamma \|\mathbf{w}\|_1 \\ &= \underset{\tilde{\mathbf{w}}}{\operatorname{argmin}} L(\mathcal{D}; \tilde{\mathbf{w}}) + \gamma \sum_{j=1}^D |w_j|; \gamma \geq 0.\end{aligned}$$

ℓ_1 -regularization

Discussion

- ℓ_1 -regularization achieves **feature selection**.

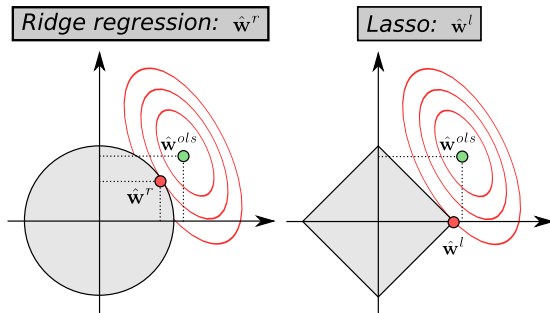


Illustration by Alexandre Gramfort, Telecom ParisTech

ℓ_1 -regularization

Discussion

- ℓ_1 -regularization achieves **feature selection**.
- The problem is still **concave**, but...
- **Difficulties**:
 - The regularizer is **not differentiable** at zero yielding **non-smooth** optimization problem.
- specific optimization techniques needed (Yuan et al., 2010).
 - In configurations with groups of highly correlated features:
 - ▶ ℓ_1 -regularization tends to select randomly one feature in each group;
 - ▶ ℓ_2 -regularization tends to yield better prediction performance.
- Consider the **elastic net** model (Hastie et al., 2009):

$$L(\mathcal{D}; \tilde{\mathbf{w}}) + \gamma_2 \|\mathbf{w}\|_2^2 + \gamma_1 \|\mathbf{w}\|_1$$

ℓ_1 -regularization

Discussion

- ℓ_1 -regularization achieves **feature selection**.
- The problem is still **concave**, but...
- **Difficulties**:
 - The regularizer is **not differentiable** at zero yielding **non-smooth** optimization problem.
 - specific optimization techniques needed (Yuan et al., 2010).
 - In configurations with groups of highly correlated features:
 - ▶ ℓ_1 -regularization tends to select randomly one feature in each group;
 - ▶ ℓ_2 -regularization tends to yield better prediction performance.
 - Consider the **elastic net** model (Hastie et al., 2009):

$$L(\mathcal{D}; \tilde{\mathbf{w}}) + \gamma_2 \|\mathbf{w}\|_2^2 + \gamma_1 \|\mathbf{w}\|_1$$

Kernel logistic regression (KLR)

- Let \mathcal{K} : positive definite kernel and $\mathcal{H}_{\mathcal{K}}$: the **RKHS** generated by \mathcal{K} .
- Let $\phi \in \mathcal{H}_{\mathcal{K}}$, a feature mapping to $\mathcal{H}_{\mathcal{K}}$.

KLR model

$$p(y_i|\mathbf{x}_i) = \frac{1}{1 + \exp -g(\mathbf{x}_i)}; \quad g(\mathbf{x}) = w_0 + \mathbf{w}^T \phi(\mathbf{x})$$

KLR model estimation problem:

$$\min_{\tilde{\mathbf{w}}} L(\tilde{\mathbf{w}}) = - \sum_{i=1}^N [y_i g(\mathbf{x}_i) - \log(1 + \exp g(\mathbf{x}_i))]$$

Kernel logistic regression (KLR)

- Let \mathcal{K} : positive definite kernel and $\mathcal{H}_{\mathcal{K}}$: the **RKHS** generated by \mathcal{K} .
- Let $\phi \in \mathcal{H}_{\mathcal{K}}$, a feature mapping to $\mathcal{H}_{\mathcal{K}}$.

KLR model

$$p(y_i|\mathbf{x}_i) = \frac{1}{1 + \exp -g(\mathbf{x}_i)}; \quad g(\mathbf{x}) = w_0 + \mathbf{w}^T \phi(\mathbf{x})$$

KLR model estimation problem:

$$\min_{\tilde{\mathbf{w}}} L(\tilde{\mathbf{w}}) = - \sum_{i=1}^N [y_i g(\mathbf{x}_i) - \log(1 + \exp g(\mathbf{x}_i))]$$

Regularized KLR

Regularized KLR model estimation problem:

$$\min_{\tilde{\mathbf{w}}} L(\tilde{\mathbf{w}}) = - \sum_{i=1}^N [y_i g(\mathbf{x}_i) - \log(1 + \exp g(\mathbf{x}_i))] + \frac{\gamma}{2} \|\mathbf{g}\|_{\mathcal{H}_{\mathcal{K}}}^2$$

- By the **representer theorem**: $g(\mathbf{x}) = w_0 + \sum_{i=1}^N \alpha_i \mathcal{K}(\mathbf{x}_i, \mathbf{x})$
- The problem is **strictly convex** and can be solved using classic solvers.

Regularized KLR

Regularized KLR model estimation problem:

$$\min_{\tilde{\mathbf{w}}} L(\tilde{\mathbf{w}}) = - \sum_{i=1}^N [y_i g(\mathbf{x}_i) - \log(1 + \exp g(\mathbf{x}_i))] + \frac{\gamma}{2} \|g\|_{\mathcal{H}_{\mathcal{K}}}^2$$

- By the **representer theorem**: $g(\mathbf{x}) = w_0 + \sum_{i=1}^N \alpha_i \mathcal{K}(\mathbf{x}_i, \mathbf{x})$
- The problem is **strictly convex** and can be solved using classic solvers.

KLR vs SVM

- It can be shown that KLR and SVM are quite related (*see Appendix*).
- Very similar prediction performance and optimal margin properties.
- Same refinements are possible: SMO, MKL...

+ Provides well-calibrated **class probabilities**.

+ Naturally generalizes to multi-class problems.

- No support vectors! → *Import Vector Machines* (Zhu and Hastie, 2002).

KLR vs SVM

- It can be shown that KLR and SVM are quite related (*see Appendix*).
 - Very similar prediction performance and optimal margin properties.
 - Same refinements are possible: SMO, MKL...
-
- + Provides well-calibrated **class probabilities**.
 - + Naturally generalizes to multi-class problems.
-
- No support vectors! → *Import Vector Machines* (Zhu and Hastie, 2002).

KLR vs SVM


- It can be shown that KLR and SVM are quite related (*see Appendix*).
 - Very similar prediction performance and optimal margin properties.
 - Same refinements are possible: SMO, MKL...
-
- + Provides well-calibrated **class probabilities**.
 - + Naturally generalizes to multi-class problems.
-
- No support vectors! → *Import Vector Machines* (Zhu and Hastie, 2002).

- ▶ Introduction
- ▶ The logistic regression model
- ▶ **Conditional Random Fields (for sequential data)**
 - Introduction
 - Inference
 - Parameter estimation
- ▶ Improvements and extensions to original CRFs
- ▶ Conclusion
- ▶ References
- ▶ Appendix

Structured-output data

John saw the big table
noun *verb* *det* *adj* *noun*

POS tagging



Musical instrument classification

Chord transcription

Recalling the notations

- \underline{x} is a sequence of observations: $\underline{x} = (x_1, \dots, x_n)$.
- \underline{y} is the corresponding sequence of labels: $\underline{y} = (y_1, \dots, y_n)$.
- We assume we have a training dataset \mathcal{D} of N (i.i.d) such sequences:
 $\mathcal{D} = \{(\underline{x}^{(1)}, \underline{y}^{(1)}), \dots, (\underline{x}^{(N)}, \underline{y}^{(N)})\}$.
- Remarks:
 - Observations are no longer assumed to be i.i.d within each sequence.
 - Sequences $\underline{x}^{(q)}$ do not necessarily have the same length, when needed n_q will denote the length of $\underline{x}^{(q)}$.

Recalling the notations

- $\underline{\mathbf{x}}$ is a sequence of observations: $\underline{\mathbf{x}} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$.
- \underline{y} is the corresponding sequence of labels: $\underline{y} = (y_1, \dots, y_n)$.
- We assume we have a training dataset \mathcal{D} of N (i.i.d) such sequences:
 $\mathcal{D} = \{(\underline{\mathbf{x}}^{(1)}, \underline{y}^{(1)}), \dots, (\underline{\mathbf{x}}^{(N)}, \underline{y}^{(N)})\}$.
- **Remarks:**
 - Observations are no longer assumed to be i.i.d within each sequence.
 - Sequences $\underline{\mathbf{x}}^{(q)}$ do not necessarily have the same length, when needed n_q will denote the length of $\underline{\mathbf{x}}^{(q)}$.

The CRF model

CRF model definition

$$\begin{aligned} p(\underline{y}|\underline{x}; \boldsymbol{\theta}) &= \frac{1}{Z(\underline{x}, \boldsymbol{\theta})} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y}) \\ &= \frac{1}{Z(\underline{x}, \boldsymbol{\theta})} \exp \Psi(\underline{x}, \underline{y}; \boldsymbol{\theta}); \quad \boldsymbol{\theta} = \{\theta_1, \dots, \theta_D\}. \end{aligned}$$

- $Z(\underline{x}, \boldsymbol{\theta}) = \sum_{\underline{y}} \exp \sum_j \theta_j F_j(\underline{x}, \underline{y})$ is called a **partition function**.
- $\Psi(\underline{x}, \underline{y}; \boldsymbol{\theta}) = \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y})$ is called a **potential function**.
- **Remarks:**
 - CRFs appear to be an extension of logistic regression to structured data.
 - Feature functions $F_j(\underline{x}, \underline{y})$ depend on the whole sequence of observations \underline{x} and labels \underline{y} .

The CRF model

CRF model definition

$$\begin{aligned} p(\underline{y}|\underline{x}; \boldsymbol{\theta}) &= \frac{1}{Z(\underline{x}, \boldsymbol{\theta})} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y}) \\ &= \frac{1}{Z(\underline{x}, \boldsymbol{\theta})} \exp \Psi(\underline{x}, \underline{y}; \boldsymbol{\theta}); \quad \boldsymbol{\theta} = \{\theta_1, \dots, \theta_D\}. \end{aligned}$$

- $Z(\underline{x}, \boldsymbol{\theta}) = \sum_{\underline{y}} \exp \sum_j \theta_j F_j(\underline{x}, \underline{y})$ is called a **partition function**.
- $\Psi(\underline{x}, \underline{y}; \boldsymbol{\theta}) = \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y})$ is called a **potential function**.
- **Remarks:**
 - CRFs appear to be an extension of logistic regression to structured data.
 - Feature functions $F_j(\underline{x}, \underline{y})$ depend on the whole sequence of observations \underline{x} and labels \underline{y} .

Defining label constraints

- Without any further assumptions on the structure of \underline{y} the model is hardly usable: one needs to enumerate all possible sequences \underline{y} for:
 - $Z(\underline{x}, \theta) = \sum_{\underline{y}} \exp \sum_j \theta_j F_j(\underline{x}, \underline{y})$;
 - $\hat{\underline{y}} = \operatorname{argmax}_{\underline{y}} p(\underline{y} | \underline{x}; \theta)$.
- Consider feature functions $F_j(\underline{x}, \underline{y})$ such that:

$$F_j(\underline{x}, \underline{y}) = \sum_{i=1}^n f_j(y_{i-1}, y_i, \underline{x}, i) ; \quad \text{where } n \text{ is the length of } \underline{x}.$$

→ defines **linear-chain** CRFs: at each position i , $1 \leq i \leq n$, each f_j depends on the **whole observation sequence**, but only on the current and previous labels.

Defining label constraints

- Without any further assumptions on the structure of \underline{y} the model is hardly usable: one needs to enumerate all possible sequences \underline{y} for:
 - $Z(\underline{x}, \theta) = \sum_{\underline{y}} \exp \sum_j \theta_j F_j(\underline{x}, \underline{y})$;
 - $\hat{\underline{y}} = \operatorname{argmax}_{\underline{y}} p(\underline{y} | \underline{x}; \theta)$.
- Consider feature functions $F_j(\underline{x}, \underline{y})$ such that:

$$F_j(\underline{x}, \underline{y}) = \sum_{i=1}^n f_j(y_{i-1}, y_i, \underline{x}, i) ; \quad \text{where } n \text{ is the length of } \underline{x}.$$

- defines **linear-chain** CRFs: at each position i , $1 \leq i \leq n$, each f_j depends on the **whole observation sequence**, but only on the current and previous labels.

Defining label constraints

$$F_j(\underline{\mathbf{x}}, \underline{y}) = \sum_{i=1}^n f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i)$$

Examples of such feature functions (for discrete observations):

- *The current observation is “**saw**”, the current label is **verb** and the previous is **noun**;*
- *The past 4 observations..., the current label is...*
- *The next observation is...*
- *The current label is...*

Defining label constraints

- For convenience, one can define two types of feature functions:
 - **Observation** (aka **state**) feature functions: $b_j(y_i, \underline{\mathbf{x}}, i)$;
 - **Transition** feature functions: $t_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i)$.

- Hence

$$p(\underline{y}|\underline{\mathbf{x}}; \boldsymbol{\theta}) = \frac{1}{Z(\underline{\mathbf{x}}, \boldsymbol{\theta})} \exp \left\{ \sum_{i=1}^n \sum_{j=1}^{D_o} \theta_j b_j(y_i, \underline{\mathbf{x}}, i) + \sum_{i=1}^n \sum_{j=1}^{D_t} \theta_j t_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i) \right\}$$

Connection to HMM

The Hidden Markov Model

$$p_{hmm}(\underline{y}, \underline{x}) \triangleq \prod_{i=1}^n p(y_i | y_{i-1}) p(\mathbf{x}_i | y_i) \quad ; \quad \text{where } p(y_1 | y_0) \triangleq p(y_1).$$

One can write:

$$\begin{aligned} p_{hmm}(\underline{y}, \underline{x}) &= \exp \left\{ \sum_{i=1}^n \log p(y_i | y_{i-1}) + \sum_{i=1}^n \log p(\mathbf{x}_i | y_i) \right\} \\ &= \exp \left\{ \sum_{i=1}^n \sum_{l, q \in \mathcal{Y}} \lambda_{lq} \mathbb{I}(y_i = l) \mathbb{I}(y_{i-1} = q) \right. \\ &\quad \left. + \sum_{i=1}^n \sum_{l \in \mathcal{Y}, \mathbf{o} \in \mathcal{X}} \mu_{\mathbf{o}l} \mathbb{I}(y_i = l) \mathbb{I}(\mathbf{x}_i = \mathbf{o}) \right\} ; \end{aligned}$$

where $\lambda_{lq} = \log p(y_i = l | y_{i-1} = q)$ and $\mu_{\mathbf{o}l} = \log p(\mathbf{x}_i = \mathbf{o} | y_i = l)$.

Connection to HMM

The Hidden Markov Model

$$p_{hmm}(\underline{y}, \underline{\mathbf{x}}) \triangleq \prod_{i=1}^n p(y_i | y_{i-1}) p(\mathbf{x}_i | y_i) \quad ; \quad \text{where } p(y_1 | y_0) \triangleq p(y_1).$$

One can write:

$$\begin{aligned} p_{hmm}(\underline{y}, \underline{\mathbf{x}}) &= \exp \left\{ \sum_{i=1}^n \log p(y_i | y_{i-1}) + \sum_{i=1}^n \log p(\mathbf{x}_i | y_i) \right\} \\ &= \exp \left\{ \sum_{i=1}^n \sum_{l, q \in \mathcal{Y}} \lambda_{lq} \mathbb{I}(y_i = l) \mathbb{I}(y_{i-1} = q) \right. \\ &\quad \left. + \sum_{i=1}^n \sum_{l \in \mathcal{Y}, \mathbf{o} \in \mathcal{X}} \mu_{\mathbf{o}l} \mathbb{I}(y_i = l) \mathbb{I}(\mathbf{x}_i = \mathbf{o}) \right\}; \end{aligned}$$

where $\lambda_{lq} = \log p(y_i = l | y_{i-1} = q)$ and $\mu_{\mathbf{o}l} = \log p(\mathbf{x}_i = \mathbf{o} | y_i = l)$.

Connection to HMM

- Using the feature functions:
 - $b_j(y_i, \underline{\mathbf{x}}, i) = \mathbb{I}(y = l)\mathbb{I}(\mathbf{x}_i = \mathbf{o})$, where each j indexes a different “ l, \mathbf{o} configuration”;
 - $t_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i) = \mathbb{I}(y_i = l)\mathbb{I}(y_{i-1} = q)$, where j indexes a different “ l, q configuration”;
- also using $p(\underline{y}|\underline{\mathbf{x}}) = \frac{p(\underline{y}, \underline{\mathbf{x}})}{\sum_{\underline{y}'} p(\underline{y}', \underline{\mathbf{x}})}$ and letting $Z(\underline{\mathbf{x}}) = \sum_{\underline{y}'} p(\underline{y}', \underline{\mathbf{x}})$, one gets:

$$p_{hmm}(\underline{y}|\underline{\mathbf{x}}) = \frac{1}{Z(\underline{\mathbf{x}})} \exp \left\{ \sum_{i=1}^n \sum_j \theta_j b_j(y_i, \underline{\mathbf{x}}, i) + \sum_{i=1}^n \sum_j \theta_j t_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i) \right\};$$

→ HMMs are a particular type of linear-chain CRFs.

Connection to HMM


- Using the feature functions:
 - $b_j(y_i, \underline{\mathbf{x}}, i) = \mathbb{I}(y = l)\mathbb{I}(\mathbf{x}_i = \mathbf{o})$, where each j indexes a different “ l, \mathbf{o} configuration”;
 - $t_j(y_{i-1}, y_i, \mathbf{x}, i) = \mathbb{I}(y_i = l)\mathbb{I}(y_{i-1} = q)$, where j indexes a different “ l, q configuration”;
- also using $p(\underline{y}|\underline{\mathbf{x}}) = \frac{p(\underline{y}, \underline{\mathbf{x}})}{\sum_{\underline{y}'} p(\underline{y}', \underline{\mathbf{x}})}$ and letting $Z(\underline{\mathbf{x}}) = \sum_{\underline{y}'} p(\underline{y}', \underline{\mathbf{x}})$, one gets:

$$p_{hmm}(\underline{y}|\underline{\mathbf{x}}) = \frac{1}{Z(\underline{\mathbf{x}})} \exp \left\{ \sum_{i=1}^n \sum_j \theta_j b_j(y_i, \underline{\mathbf{x}}, i) + \sum_{i=1}^n \sum_j \theta_j t_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i) \right\};$$


→ HMMs are a particular type of linear-chain CRFs.

Connection to HMM

Discussion



HMM



CRF

CRFs have a number of advantages over HMMs, as a consequence of two major differences:

- CRFs are **discriminative models**.
- CRFs are **undirected models**.

Connection to HMM

Advantage of the discriminative nature of CRF

HMM: observation x_i is independent of all other variables given its parent state y_i .

CRF: no assumptions on the dependencies among the observations: only $p(\underline{y}|\underline{x})$ is modeled.

→ CRFs can safely:

- exploit **overlapping features**;
- account for **long-term dependencies**, considering the whole sequence of observations \underline{x} at each location i ($i \mapsto b_j(y_i, \underline{x}, i)$);
- use **transition feature-functions** $t_j(y_{i-1}, y_i, \underline{x}, i)$.

Using linear-chain CRFs

- **Inference:** given a model θ , how to compute:
 - $\hat{y} = \operatorname{argmax}_{\underline{y}} p(\underline{y}|\underline{x}; \theta)$?
 - $Z(\underline{x}, \theta) = \sum_{\underline{y}} \exp \sum_j F_j(\underline{x}, \underline{y})$ to deduce
 $p(\underline{y}|\underline{x}; \theta) = \frac{1}{Z(\underline{x}, \theta)} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y})$?
- **Parameter estimation:** given a training dataset
 $\mathcal{D} = \{(\underline{x}^{(1)}, \underline{y}^{(1)}), \dots, (\underline{x}^{(N)}, \underline{y}^{(N)})\}$, how to estimate the optimal θ ?

Using linear-chain CRFs

- **Inference:** given a model θ , how to compute:
 - $\hat{y} = \operatorname{argmax}_{\underline{y}} p(\underline{y}|\underline{x}; \theta)$?
 - $Z(\underline{x}, \theta) = \sum_{\underline{y}} \exp \sum_j F_j(\underline{x}, \underline{y})$ to deduce
 $p(\underline{y}|\underline{x}; \theta) = \frac{1}{Z(\underline{x}, \theta)} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, \underline{y})$?
- **Parameter estimation:** given a training dataset
 $\mathcal{D} = \{(\underline{x}^{(1)}, \underline{y}^{(1)}), \dots, (\underline{x}^{(N)}, \underline{y}^{(N)})\}$, how to estimate the optimal θ ?

Decoding the optimal sequence

- **Problem:** solve $\hat{\underline{y}} = \operatorname{argmax}_{\underline{y} \in \mathcal{Y}^n} p(\underline{y} | \underline{\mathbf{x}}; \boldsymbol{\theta})$, with $|\mathcal{Y}|^n$ possible assignments!
- **Solution:** use the **Viterbi** algorithm.
- **Remark:**

$$\begin{aligned}\hat{\underline{y}} = \operatorname{argmax}_{\underline{y}} p(\underline{y} | \underline{\mathbf{x}}; \boldsymbol{\theta}) &= \operatorname{argmax}_{\underline{y}} \frac{1}{Z(\underline{\mathbf{x}}, \boldsymbol{\theta})} \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y}) \\ &= \operatorname{argmax}_{\underline{y}} \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y}) \\ &= \operatorname{argmax}_{\underline{y}} \sum_{i=1}^n \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i)\end{aligned}$$

Decoding the optimal sequence

- **Problem:** solve $\hat{\underline{y}} = \operatorname{argmax}_{\underline{y} \in \mathcal{Y}^n} p(\underline{y} | \underline{\mathbf{x}}; \boldsymbol{\theta})$, with $|\mathcal{Y}|^n$ possible assignments!
- **Solution:** use the **Viterbi** algorithm.
- **Remark:**

$$\begin{aligned}\hat{\underline{y}} = \operatorname{argmax}_{\underline{y}} p(\underline{y} | \underline{\mathbf{x}}; \boldsymbol{\theta}) &= \operatorname{argmax}_{\underline{y}} \frac{1}{Z(\underline{\mathbf{x}}, \boldsymbol{\theta})} \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y}) \\ &= \operatorname{argmax}_{\underline{y}} \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y}) \\ &= \operatorname{argmax}_{\underline{y}} \sum_{i=1}^n \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i)\end{aligned}$$

Decoding the optimal sequence

- **Problem:** solve $\hat{\underline{y}} = \operatorname{argmax}_{\underline{y} \in \mathcal{Y}^n} p(\underline{y}|\underline{\mathbf{x}}; \boldsymbol{\theta})$, with $|\mathcal{Y}|^n$ possible assignments!
- **Solution:** use the **Viterbi** algorithm.
- **Remark:**

$$\begin{aligned}\hat{\underline{y}} = \operatorname{argmax}_{\underline{y}} p(\underline{y}|\underline{\mathbf{x}}; \boldsymbol{\theta}) &= \operatorname{argmax}_{\underline{y}} \frac{1}{Z(\underline{\mathbf{x}}, \boldsymbol{\theta})} \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y}) \\ &= \operatorname{argmax}_{\underline{y}} \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y}) \\ &= \operatorname{argmax}_{\underline{y}} \sum_{i=1}^n \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i)\end{aligned}$$

Decoding the optimal sequence

Let: $g_i(y_{i-1}, y_i) \triangleq \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{x}, i)$; then:

$$\hat{y} = \underset{\underline{y}}{\operatorname{argmax}} \sum_{i=1}^n \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{x}, i) = \underset{\underline{y}}{\operatorname{argmax}} \sum_{i=1}^n g_i(y_{i-1}, y_i).$$

Let $\delta_m(s)$ be the optimal “intermediate score” where at time step m the label value is s :

$$\delta_m(s) \triangleq \max_{\{y_1, \dots, y_{m-1}\}} \left[\sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + g_m(y_{m-1}, s) \right]$$

Decoding the optimal sequence

Let: $g_i(y_{i-1}, y_i) \triangleq \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{x}, i)$; then:

$$\hat{y} = \underset{\underline{y}}{\operatorname{argmax}} \sum_{i=1}^n \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{x}, i) = \underset{\underline{y}}{\operatorname{argmax}} \sum_{i=1}^n g_i(y_{i-1}, y_i).$$


Let $\delta_m(s)$ be the optimal “intermediate score” where at time step m the label value is s :

$$\delta_m(s) \triangleq \max_{\{y_1, \dots, y_{m-1}\}} \left[\sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + g_m(y_{m-1}, s) \right]$$

Decoding the optimal sequence

Trellis representation

$$\delta_m(\mathbf{s}) \triangleq \max_{\{y_1, \dots, y_{m-1}\}} \left[\sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + g_m(y_{m-1}, \mathbf{s}) \right]$$



Decoding the optimal sequence

- The intermediate scores $\delta_m(s)$ can be efficiently computed using¹:

Viterbi recursion

$$\delta_m(s) = \max_{y_{m-1} \in \mathcal{Y}} [\delta_{m-1}(y_{m-1}) + g_m(y_{m-1}, s)]; 1 \leq m \leq n$$

- As we proceed we need to keep track of the selected predecessor of s , at each time step m .
- We use $\psi_m(s)$ for this purpose.

¹See *Appendix for more details.*

Decoding the optimal sequence

Viterbi algorithm

Initialization:

$$\delta_1(s) = g_1(y_0, s); \forall s \in \mathcal{Y}; y_0 = \text{start}$$

$$\psi_1(s) = \text{start}$$

Recursion:

$$\forall s \in \mathcal{Y}; 1 \leq m \leq n$$

$$\delta_m(s) = \max_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

$$\psi_m(s) = \operatorname{argmax}_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

Termination:

$$\delta_n(y_n^*) = \max_{y \in \mathcal{Y}} \delta_n(y)$$

$$y_n^* = \operatorname{argmax}_{y \in \mathcal{Y}} \delta_n(y)$$

Path backtracking:

$$y_m^* = \psi_{m+1}(y_{m+1}^*); m = n-1, n-2, \dots, 1.$$

Decoding the optimal sequence

Viterbi algorithm

Initialization:

$$\delta_1(s) = g_1(y_0, s); \forall s \in \mathcal{Y}; y_0 = \text{start}$$

$$\psi_1(s) = \text{start}$$

Recursion:

$$\forall s \in \mathcal{Y}; 1 \leq m \leq n$$

$$\delta_m(s) = \max_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

$$\psi_m(s) = \operatorname{argmax}_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

Termination:

$$\delta_n(y_n^*) = \max_{y \in \mathcal{Y}} \delta_n(y)$$

$$y_n^* = \operatorname{argmax}_{y \in \mathcal{Y}} \delta_n(y)$$

Path backtracking:

$$y_m^* = \psi_{m+1}(y_{m+1}^*); m = n-1, n-2, \dots, 1.$$

Decoding the optimal sequence

Viterbi algorithm

Initialization:

$$\delta_1(s) = g_1(y_0, s); \forall s \in \mathcal{Y}; y_0 = \text{start}$$

$$\psi_1(s) = \text{start}$$

Recursion:

$$\forall s \in \mathcal{Y}; 1 \leq m \leq n$$

$$\delta_m(s) = \max_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

$$\psi_m(s) = \operatorname{argmax}_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

Termination:

$$\delta_n(y_n^*) = \max_{y \in \mathcal{Y}} \delta_n(y)$$

$$y_n^* = \operatorname{argmax}_{y \in \mathcal{Y}} \delta_n(y)$$

Path backtracking:

$$y_m^* = \psi_{m+1}(y_{m+1}^*); m = n-1, n-2, \dots, 1.$$

Decoding the optimal sequence

Viterbi algorithm

Initialization:

$$\delta_1(s) = g_1(y_0, s); \forall s \in \mathcal{Y}; y_0 = \text{start}$$

$$\psi_1(s) = \text{start}$$

Recursion:

$$\forall s \in \mathcal{Y}; 1 \leq m \leq n$$

$$\delta_m(s) = \max_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

$$\psi_m(s) = \operatorname{argmax}_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

Termination:

$$\delta_n(y_n^*) = \max_{y \in \mathcal{Y}} \delta_n(y)$$

$$y_n^* = \operatorname{argmax}_{y \in \mathcal{Y}} \delta_n(y)$$


Path backtracking:

$$y_m^* = \psi_{m+1}(y_{m+1}^*); m = n-1, n-2, \dots, 1.$$

Decoding the optimal sequence

Backtracking

$$y_m^* = \psi_{m+1}(y_{m+1}^*); \quad m = n-1, n-2, \dots, 1.$$



Complexity of Viterbi decoding

Remarks on the computational cost:

- $O(K^2n)$ in the worst case; $K = |\mathcal{Y}|$.
- In practice: $O(\mathcal{T}Kn)$, where \mathcal{T} : average number of possible “transitions” between labels y .
- Can be reduced using **beam search**: exploring a subset of possible labels at each time position (Ortmanns et al., 1996).

Complexity of Viterbi decoding

Remarks on the computational cost:

- $O(K^2n)$ in the worst case; $K = |\mathcal{Y}|$.
- In practice: $O(\mathcal{T}Kn)$, where \mathcal{T} : average number of possible “transitions” between labels y .
- Can be reduced using **beam search**: exploring a subset of possible labels at each time position (Ortmanns et al., 1996).

Computing the partition function $Z(\underline{\mathbf{x}}, \boldsymbol{\theta})$

The sum-product problem

Recall the CRF model:

$$p(\underline{y}|\underline{\mathbf{x}}; \boldsymbol{\theta}) = \frac{1}{Z(\underline{\mathbf{x}}, \boldsymbol{\theta})} \prod_{i=1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}});$$

$$M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) = \exp \left(\sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i) \right);$$

$$Z(\underline{\mathbf{x}}, \boldsymbol{\theta}) = \sum_{\underline{y} \in \mathcal{Y}^n} \prod_{i=1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) : \text{intractable as is...}$$

→ use the **forward-backward** method: reduces **complexity** from $O(K^n)$ to $O(nK^2)$.

The forward-backward method

- Defining $\alpha_m(y_m) = \sum_{y_{m-1}} M_m(y_{m-1}, y_m) \alpha_{m-1}(y_{m-1})$; $2 \leq m \leq n$, it is easily shown² that:

At the end of the sequence

$$Z(\underline{\mathbf{x}}, \theta) = \sum_{y_n \in \mathcal{Y}} \alpha_n(y_n).$$

- Alternatively, defining $\beta_m(y_m) = \sum_{y_{m+1}} M_{m+1}(y_m, y_{m+1}) \beta_{m+1}(y_{m+1})$; $1 \leq m \leq n-1$ and $\beta_n(y_n) = 1$, one gets:

At the beginning of the sequence

$$Z(\underline{\mathbf{x}}, \theta) = \sum_{y_1 \in \mathcal{Y}} M_1(y_0, y_1) \beta_1(y_1).$$

²See Appendix for more details

The forward-backward method

- Defining $\alpha_m(y_m) = \sum_{y_{m-1}} M_m(y_{m-1}, y_m) \alpha_{m-1}(y_{m-1})$; $2 \leq m \leq n$, it is easily shown² that:

At the end of the sequence

$$Z(\underline{\mathbf{x}}, \theta) = \sum_{y_n \in \mathcal{Y}} \alpha_n(y_n).$$

- Alternatively, defining $\beta_m(y_m) = \sum_{y_{m+1}} M_{m+1}(y_m, y_{m+1}) \beta_{m+1}(y_{m+1})$; $1 \leq m \leq n-1$ and $\beta_n(y_n) = 1$, one gets:

At the beginning of the sequence

$$Z(\underline{\mathbf{x}}, \theta) = \sum_{y_1 \in \mathcal{Y}} M_1(y_0, y_1) \beta_1(y_1).$$

²See Appendix for more details

Marginal probability

$$p(y_{m-1}, y_m | \underline{\mathbf{x}}) = \sum_{\underline{y} \setminus \{y_{m-1}, y_m\}} p(\underline{y} | \mathbf{x})$$

Marginal probability by forward-backward

$$p(y_{m-1}, y_m | \underline{\mathbf{x}}) = \frac{1}{Z(\underline{\mathbf{x}})} \alpha_{m-1}(y_{m-1}) M_m(y_{m-1}, y_m, \underline{\mathbf{x}}) \beta_m(y_m).$$

More details in the appendix.

Marginal probability

$$p(y_{m-1}, y_m | \underline{\mathbf{x}}) = \sum_{\underline{y} \setminus \{y_{m-1}, y_m\}} p(\underline{y} | \mathbf{x})$$

Marginal probability by forward-backward

$$p(y_{m-1}, y_m | \underline{\mathbf{x}}) = \frac{1}{Z(\underline{\mathbf{x}})} \alpha_{m-1}(y_{m-1}) M_m(y_{m-1}, y_m, \underline{\mathbf{x}}) \beta_m(y_m).$$

More details in the appendix.

- ▶ Introduction
- ▶ The logistic regression model
- ▶ Conditional Random Fields (for sequential data)
 - Introduction
 - Inference
 - Parameter estimation
- ▶ Improvements and extensions to original CRFs
- ▶ Conclusion
- ▶ References
- ▶ Appendix

Negative log-likelihood (NLL)

- Given training data $\mathcal{D} = \{(\underline{\mathbf{x}}^{(1)}, \underline{y}^{(1)}), \dots, (\underline{\mathbf{x}}^{(N)}, \underline{y}^{(N)})\}$, the NLL is:

$$\begin{aligned}
 L(\mathcal{D}; \theta) &= - \sum_{q=1}^N \log p(\underline{y}^{(q)} | \underline{\mathbf{x}}^{(q)}; \theta) \\
 &= \sum_{q=1}^N \left\{ \log Z(\underline{\mathbf{x}}^{(q)}; \theta) - \sum_{i=1}^{n_q} \sum_{j=1}^D \theta_j f_j(y_{i-1}^{(q)}, y_i^{(q)}, \underline{\mathbf{x}}^{(q)}, i) \right\} \\
 &= \sum_{q=1}^N \left\{ \log Z(\underline{\mathbf{x}}^{(q)}; \theta) - \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}^{(q)}, \underline{y}^{(q)}) \right\}.
 \end{aligned}$$

- $L(\mathcal{D}; \theta)$ is **convex** \rightarrow gradient-descent will converge to global minimum.

NLL gradient

Gradient:
$$\frac{\partial L(\mathcal{D}; \boldsymbol{\theta})}{\partial \theta_k} = \sum_{q=1}^N \left\{ \frac{\partial}{\partial \theta_k} \log Z(\underline{\mathbf{x}}^{(q)}; \boldsymbol{\theta}) - F_k(\underline{\mathbf{x}}^{(q)}, \underline{y}^{(q)}) \right\}.$$

$$\begin{aligned} \frac{\partial}{\partial \theta_k} \log Z(\underline{\mathbf{x}}; \boldsymbol{\theta}) &= \frac{1}{Z(\underline{\mathbf{x}}; \boldsymbol{\theta})} \sum_{\underline{y} \in \mathcal{Y}^n} \frac{\partial}{\partial \theta_k} \left[\exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y}) \right] \\ &= \frac{1}{Z(\underline{\mathbf{x}}; \boldsymbol{\theta})} \sum_{\underline{y} \in \mathcal{Y}^n} F_k(\underline{\mathbf{x}}, \underline{y}) \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, \underline{y}) \\ &= \sum_{\underline{y} \in \mathcal{Y}^n} F_k(\underline{\mathbf{x}}, \underline{y}) \frac{\exp \sum_j \theta_j F_j(\underline{\mathbf{x}}, \underline{y})}{Z(\underline{\mathbf{x}}; \boldsymbol{\theta})} \\ &= \sum_{\underline{y} \in \mathcal{Y}^n} F_k(\underline{\mathbf{x}}, \underline{y}) p(\underline{y} | \underline{\mathbf{x}}; \boldsymbol{\theta}) \\ &= \mathbb{E}_{p(\underline{y} | \underline{\mathbf{x}}; \boldsymbol{\theta})} [F_k(\underline{\mathbf{x}}, \underline{y})]. \end{aligned}$$

NLL gradient

Gradient:
$$\frac{\partial L(\mathcal{D}; \boldsymbol{\theta})}{\partial \theta_k} = \sum_{q=1}^N \left\{ \frac{\partial}{\partial \theta_k} \log Z(\underline{\mathbf{x}}^{(q)}; \boldsymbol{\theta}) - F_k(\underline{\mathbf{x}}^{(q)}, \underline{y}^{(q)}) \right\}.$$

$\frac{\partial}{\partial \theta_k} \log Z(\underline{\mathbf{x}}; \boldsymbol{\theta}) = \mathbb{E}_{p(\underline{y}|\underline{\mathbf{x}};\boldsymbol{\theta})} [F_k(\underline{\mathbf{x}}, \underline{y})]$: conditional expectation given $\underline{\mathbf{x}}$.

$$\frac{\partial L(\mathcal{D}; \boldsymbol{\theta})}{\partial \theta_k} = \sum_{q=1}^N \left\{ \mathbb{E}_{p(\underline{y}|\underline{\mathbf{x}}^{(q)};\boldsymbol{\theta})} [F_k(\underline{\mathbf{x}}^{(q)}, \underline{y})] - F_k(\underline{\mathbf{x}}^{(q)}, \underline{y}^{(q)}) \right\}.$$

Optimality condition

- Setting the derivatives to 0, i.e. $\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = 0$, yields:

$$\sum_{q=1}^N \mathbb{E}_{p(\underline{y}|\underline{\mathbf{x}}^{(q)}; \theta)} [F_k(\underline{\mathbf{x}}^{(q)}, \underline{y})] = \sum_{q=1}^N F_k(\underline{\mathbf{x}}^{(q)}, \underline{y}^{(q)}); \quad 1 \leq k \leq D$$

- No closed-form solution: numerical optimization is again needed.
- Need to compute $\mathbb{E}_{p(\underline{y}|\underline{\mathbf{x}}^{(q)}; \theta)} [F_k(\underline{\mathbf{x}}^{(q)}, \underline{y})]$ efficiently.

Optimality condition

- Setting the derivatives to 0, *i.e.* $\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = 0$, yields:

$$\frac{1}{N} \sum_{q=1}^N \mathbb{E}_{p(\underline{y}|\underline{x}^{(q)}; \theta)} \left[F_k(\underline{x}^{(q)}, \underline{y}) \right] = \frac{1}{N} \sum_{q=1}^N F_k(\underline{x}^{(q)}, \underline{y}^{(q)}); \quad 1 \leq k \leq D$$

- Average expectation under the model = empirical mean.
- No closed-form solution: numerical optimization is again needed.
- Need to compute $\mathbb{E}_{p(\underline{y}|\underline{x}^{(q)}; \theta)} \left[F_k(\underline{x}^{(q)}, \underline{y}) \right]$ efficiently.

Optimality condition

- Setting the derivatives to 0, *i.e.* $\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = 0$, yields:

$$\frac{1}{N} \sum_{q=1}^N \mathbb{E}_{p(\underline{y}|\underline{x}^{(q)}; \theta)} \left[F_k(\underline{x}^{(q)}, \underline{y}) \right] = \frac{1}{N} \sum_{q=1}^N F_k(\underline{x}^{(q)}, \underline{y}^{(q)}); \quad 1 \leq k \leq D$$

- Average expectation under the model = empirical mean.
- No closed-form solution: numerical optimization is again needed.
- Need to compute $\mathbb{E}_{p(\underline{y}|\underline{x}^{(q)}; \theta)} \left[F_k(\underline{x}^{(q)}, \underline{y}) \right]$ efficiently.

Optimality condition

- Setting the derivatives to 0, *i.e.* $\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = 0$, yields:

$$\frac{1}{N} \sum_{q=1}^N \mathbb{E}_{p(\underline{y}|\underline{\mathbf{x}}^{(q)}; \theta)} \left[F_k(\underline{\mathbf{x}}^{(q)}, \underline{y}) \right] = \frac{1}{N} \sum_{q=1}^N F_k(\underline{\mathbf{x}}^{(q)}, \underline{y}^{(q)}); \quad 1 \leq k \leq D$$

- Average expectation under the model = empirical mean.
- No closed-form solution: numerical optimization is again needed.
- Need to compute $\mathbb{E}_{p(\underline{y}|\underline{\mathbf{x}}^{(q)}; \theta)} \left[F_k(\underline{\mathbf{x}}^{(q)}, \underline{y}) \right]$ efficiently.

Efficient gradient computation

$$\begin{aligned}
 \mathbb{E}_{p(\underline{y}|\underline{x};\theta)} [F_k(\underline{x}, \underline{y})] &= \sum_{\underline{y} \in \mathcal{Y}^n} F_k(\underline{x}, \underline{y}) p(\underline{y}|\underline{x}; \theta) \\
 &= \sum_{i=1}^n \sum_{\underline{y} \in \mathcal{Y}^n} f_k(y_{i-1}, y_i, \underline{x}) p(\underline{y}|\underline{x}; \theta) \\
 &= \sum_{i=1}^n \sum_{y_{i-1}, y_i \in \mathcal{Y}^2} f_k(y_{i-1}, y_i, \underline{x}) p(y_{i-1}, y_i|\underline{x}; \theta)
 \end{aligned}$$

$p(y_{i-1}, y_i|\underline{x}; \theta)$ is the **marginal probability** which thanks to the **forward-backward** algorithm is obtained by:

$$p(y_{i-1}, y_i|\underline{x}) = \frac{1}{Z(\underline{x})} \alpha_{i-1}(y_{i-1}) M_i(y_{i-1}, y_i, \underline{x}) \beta_i(y_i).$$

Efficient gradient computation

$$\begin{aligned}
 \mathbb{E}_{p(\underline{y}|\underline{x},\theta)} [F_k(\underline{x}, \underline{y})] &= \sum_{\underline{y} \in \mathcal{Y}^n} F_k(\underline{x}, \underline{y}) p(\underline{y}|\underline{x}; \theta) \\
 &= \sum_{i=1}^n \sum_{\underline{y} \in \mathcal{Y}^n} f_k(y_{i-1}, y_i, \underline{x}) p(\underline{y}|\underline{x}; \theta) \\
 &= \sum_{i=1}^n \sum_{y_{i-1}, y_i \in \mathcal{Y}^2} f_k(y_{i-1}, y_i, \underline{x}) p(y_{i-1}, y_i | \underline{x}; \theta)
 \end{aligned}$$

$p(y_{i-1}, y_i | \underline{x}; \theta)$ is the **marginal probability** which thanks to the **forward-backward** algorithm is obtained by:

$$p(y_{i-1}, y_i | \underline{x}) = \frac{1}{Z(\underline{x})} \alpha_{i-1}(y_{i-1}) M_i(y_{i-1}, y_i, \underline{x}) \beta_i(y_i).$$

Optimization

Many **algorithms** are available (see Sokolovska, 2010; Lavergne et al., 2010):

- Generalized iterative scaling (Lafferty et al., 2001): original algorithm, slow convergence, suboptimal.
- Conjugate gradient (Wallach, 2002): faster convergence, better quality.
- **L-BFGS** (McCallum, 2002): fast convergence, scalable; a good option, most used.
- Stochastic gradient: suboptimal, simple, online, **large-scale applications**.

- ▶ Introduction
- ▶ The logistic regression model
- ▶ Conditional Random Fields (for sequential data)
- ▶ Improvements and extensions to original CRFs
 - Regularization
 - Hidden-state CRF model
 - Other extensions
- ▶ Conclusion
- ▶ References
- ▶ Appendix

Regularization

Using ℓ_2 -norm

- Redefine the objective function as: $L(\mathcal{D}; \theta) = L(\mathcal{D}; \theta) + \frac{\|\theta\|_2^2}{2\sigma^2}$;
 σ^2 : a free parameter penalizing large weights (as in **ridge regression**).
- The gradient coefficients become:

$$\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = \sum_{q=1}^N \left\{ \mathbb{E}_{p(\underline{y}|\underline{x}^{(q)}; \theta)} [F_k(\underline{x}^{(q)}, \underline{y})] - F_k(\underline{x}^{(q)}, \underline{y}^{(q)}) \right\} + \frac{\theta_k}{\sigma^2}.$$
- **Advantages:**
 - The objective becomes **strictly convex**.
 - Shrinkage of θ coefficients is achieved avoiding overfitting and numerical problems.
- σ^2 needs to be tuned (usually by cross-validation).

Regularization

Using ℓ_2 -norm

- Redefine the objective function as: $L(\mathcal{D}; \theta) = L(\mathcal{D}; \theta) + \frac{\|\theta\|_2^2}{2\sigma^2}$;
 σ^2 : a free parameter penalizing large weights (as in **ridge regression**).
- The gradient coefficients become:

$$\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = \sum_{q=1}^N \left\{ \mathbb{E}_{p(\underline{y}|\underline{x}^{(q)}; \theta)} [F_k(\underline{x}^{(q)}, \underline{y})] - F_k(\underline{x}^{(q)}, \underline{y}^{(q)}) \right\} + \frac{\theta_k}{\sigma^2}.$$
- **Advantages:**
 - The objective becomes **strictly convex**.
 - Shrinkage of θ coefficients is achieved avoiding overfitting and numerical problems.
- σ^2 needs to be tuned (usually by cross-validation).

Regularization

Using ℓ_2 -norm

- Redefine the objective function as: $L(\mathcal{D}; \theta) = L(\mathcal{D}; \theta) + \frac{\|\theta\|_2^2}{2\sigma^2}$;
 σ^2 : a free parameter penalizing large weights (as in **ridge regression**).
- The gradient coefficients become:

$$\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = \sum_{q=1}^N \left\{ \mathbb{E}_{p(\underline{y}|\underline{\mathbf{x}}^{(q)}; \theta)} [F_k(\underline{\mathbf{x}}^{(q)}, \underline{y})] - F_k(\underline{\mathbf{x}}^{(q)}, \underline{y}^{(q)}) \right\} + \frac{\theta_k}{\sigma^2}.$$
- **Advantages:**
 - The objective becomes **strictly convex**.
 - Shrinkage of θ coefficients is achieved avoiding overfitting and numerical problems.
- σ^2 needs to be tuned (usually by cross-validation).

Regularization

Using ℓ_1 -norm to perform feature selection

- Redefine the objective function as:
$$L(\mathcal{D}; \theta) = L(\mathcal{D}; \theta) + \rho \|\theta\|_1 = L(\mathcal{D}; \theta) + \rho \sum_{j=1}^D |\theta_j|$$
 (as in the LASSO).
- Advantage: performs feature selection
in some NLP apps: up to 95% of the features can be discarded without affecting performance! (see Sokolovska, 2010).
- Difficulties:
 - The regularizer is **not differentiable** at zero: specific optimization techniques needed (Sokolovska, 2010).
 - In configurations with groups of highly correlated features, tend to select randomly one feature in each group.

Regularization

Using ℓ_1 -norm to perform feature selection

- Redefine the objective function as:
$$L(\mathcal{D}; \theta) = L(\mathcal{D}; \theta) + \rho \|\theta\|_1 = L(\mathcal{D}; \theta) + \rho \sum_{j=1}^D |\theta_j|$$
 (as in the LASSO).
- **Advantage:** performs **feature selection**
in some NLP apps: up to 95% of the features can be discarded without affecting performance! (see Sokolovska, 2010).
- **Difficulties:**
 - The regularizer is **not differentiable** at zero: specific optimization techniques needed (Sokolovska, 2010).
 - In configurations with groups of highly correlated features, tend to select randomly one feature in each group.


- ▶ Introduction
- ▶ The logistic regression model
- ▶ Conditional Random Fields (for sequential data)
- ▶ Improvements and extensions to original CRFs
 - Regularization
 - Hidden-state CRF model
 - Other extensions
- ▶ Conclusion
- ▶ References
- ▶ Appendix

Motivation

Problem: the CRF model does not support **hidden states**.


CRF

$$\mathcal{D} = \{(\underline{\mathbf{x}}^{(i)}, \underline{y}^{(i)})\}_i$$



Hidden-state CRF

$$\mathcal{D} = \{(\underline{\mathbf{x}}^{(i)}, y^{(i)})\}_i$$




Motivation

Problem: the CRF model does not support **hidden states**.


CRF

$$\mathcal{D} = \{(\underline{\mathbf{x}}^{(i)}, \underline{y}^{(i)})\}_i$$



Hidden-state CRF

$$\mathcal{D} = \{(\underline{\mathbf{x}}^{(i)}, y^{(i)})\}_i$$



The HCRF model

(Quattoni et al., 2007)

- Each sequence of observations $\underline{\mathbf{x}} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$ is associated with:
 - a unique label y ;
 - a sequence of **latent variables** $\underline{h} = (h_1, \dots, h_n)$, where $h_i \in \mathcal{H}$.

HCRF model definition

$$p(y, \underline{h} | \underline{\mathbf{x}}; \boldsymbol{\theta}) = \frac{1}{Z(\underline{\mathbf{x}}, \boldsymbol{\theta})} \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, y, \underline{h})$$

$$Z(\underline{\mathbf{x}}, \boldsymbol{\theta}) = \sum_{y, \underline{h}} \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}, y, \underline{h}); \quad \boldsymbol{\theta} = \{\theta_1, \dots, \theta_D\}.$$

Inference in HCRF

- HCRF model: $p(y, \underline{h} | \underline{x}; \theta) = \frac{1}{Z(\underline{x}, \theta)} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, y, \underline{h})$.
- Let $Z'(y, \underline{x}, \theta) \triangleq \sum_{\underline{h} \in \mathcal{H}^n} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, y, \underline{h})$: **marginalization** wrt \underline{h} .
- We have:
 - $Z(\underline{x}, \theta) = \sum_y Z'(y, \underline{x}, \theta)$;
 - $p(y | \underline{x}; \theta) = \sum_{\underline{h} \in \mathcal{H}^n} p(y, \underline{h} | \underline{x}; \theta) = \frac{Z'(y, \underline{x}, \theta)}{\sum_y Z'(y, \underline{x}, \theta)}$.
- $Z'(y, \underline{x}, \theta)$ can be easily computed using forward/backward recursions (as done in CRF).
- To classify new test cases, use:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} p(y | \underline{x}; \theta^*) = \operatorname{argmax}_{y \in \mathcal{Y}} \frac{Z'(y, \underline{x}, \theta)}{\sum_y Z'(y, \underline{x}, \theta)}.$$
- To compute the partition function, use $Z(\underline{x}, \theta) = \sum_y Z'(y, \underline{x}, \theta)$.

Inference in HCRF

- HCRF model: $p(y, \underline{h} | \underline{x}; \theta) = \frac{1}{Z(\underline{x}, \theta)} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, y, \underline{h})$.
- Let $Z'(y, \underline{x}, \theta) \triangleq \sum_{\underline{h} \in \mathcal{H}^n} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, y, \underline{h})$: **marginalization** wrt \underline{h} .
- We have:
 - $Z(\underline{x}, \theta) = \sum_y Z'(y, \underline{x}, \theta)$;
 - $p(y | \underline{x}; \theta) = \sum_{\underline{h} \in \mathcal{H}^n} p(y, \underline{h} | \underline{x}; \theta) = \frac{Z'(y, \underline{x}, \theta)}{\sum_y Z'(y, \underline{x}, \theta)}$.
- $Z'(y, \underline{x}, \theta)$ can be easily computed using forward/backward recursions (as done in CRF).
- To classify new test cases, use:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} p(y | \underline{x}; \theta^*) = \operatorname{argmax}_{y \in \mathcal{Y}} \frac{Z'(y, \underline{x}, \theta)}{\sum_y Z'(y, \underline{x}, \theta)}.$$
- To compute the partition function, use $Z(\underline{x}, \theta) = \sum_y Z'(y, \underline{x}, \theta)$.

Inference in HCRF

- HCRF model: $p(y, \underline{h} | \underline{x}; \theta) = \frac{1}{Z(\underline{x}, \theta)} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, y, \underline{h})$.
- Let $Z'(y, \underline{x}, \theta) \triangleq \sum_{\underline{h} \in \mathcal{H}^n} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, y, \underline{h})$: **marginalization** wrt \underline{h} .
- We have:
 - $Z(\underline{x}, \theta) = \sum_y Z'(y, \underline{x}, \theta)$;
 - $p(y | \underline{x}; \theta) = \sum_{\underline{h} \in \mathcal{H}^n} p(y, \underline{h} | \underline{x}; \theta) = \frac{Z'(y, \underline{x}, \theta)}{\sum_y Z'(y, \underline{x}, \theta)}$.
- $Z'(y, \underline{x}, \theta)$ can be easily computed using forward/backward recursions (as done in CRF).
- To classify new test cases, use:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} p(y | \underline{x}; \theta^*) = \operatorname{argmax}_{y \in \mathcal{Y}} \frac{Z'(y, \underline{x}, \theta)}{\sum_y Z'(y, \underline{x}, \theta)}.$$
- To compute the partition function, use $Z(\underline{x}, \theta) = \sum_y Z'(y, \underline{x}, \theta)$.

Inference in HCRF

- HCRF model: $p(y, \underline{h} | \underline{x}; \theta) = \frac{1}{Z(\underline{x}, \theta)} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, y, \underline{h})$.
- Let $Z'(y, \underline{x}, \theta) \triangleq \sum_{\underline{h} \in \mathcal{H}^n} \exp \sum_{j=1}^D \theta_j F_j(\underline{x}, y, \underline{h})$: **marginalization** wrt \underline{h} .
- We have:
 - $Z(\underline{x}, \theta) = \sum_y Z'(y, \underline{x}, \theta)$;
 - $p(y | \underline{x}; \theta) = \sum_{\underline{h} \in \mathcal{H}^n} p(y, \underline{h} | \underline{x}; \theta) = \frac{Z'(y, \underline{x}, \theta)}{\sum_y Z'(y, \underline{x}, \theta)}$.
- $Z'(y, \underline{x}, \theta)$ can be easily computed using forward/backward recursions (as done in CRF).
- To classify new test cases, use:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} p(y | \underline{x}; \theta^*) = \operatorname{argmax}_{y \in \mathcal{Y}} \frac{Z'(y, \underline{x}, \theta)}{\sum_y Z'(y, \underline{x}, \theta)}.$$
- To compute the partition function, use $Z(\underline{x}, \theta) = \sum_y Z'(y, \underline{x}, \theta)$.

Negative log-likelihood

$$\begin{aligned} L(\mathcal{D}; \theta) &= - \sum_{q=1}^N \log p(y^{(q)} | \underline{\mathbf{x}}^{(q)}; \theta) \\ &= \sum_{q=1}^N \left\{ \log \left(\sum_y Z'(y, \underline{\mathbf{x}}^{(q)}, \theta) \right) \right. \\ &\quad \left. - \log \left(\sum_{\underline{h} \in \mathcal{H}} \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}^{(q)}, y^{(q)}, \underline{h}) \right) \right\} \end{aligned}$$

$L(\mathcal{D}; \theta)$ is no longer convex \rightarrow convergence to a local minimum.

Negative log-likelihood

$$\begin{aligned} L(\mathcal{D}; \theta) &= - \sum_{q=1}^N \log p(y^{(q)} | \underline{\mathbf{x}}^{(q)}; \theta) \\ &= \sum_{q=1}^N \left\{ \log \left(\sum_y Z'(y, \underline{\mathbf{x}}^{(q)}, \theta) \right) \right. \\ &\quad \left. - \log \left(\sum_{\underline{h} \in \mathcal{H}} \exp \sum_{j=1}^D \theta_j F_j(\underline{\mathbf{x}}^{(q)}, y^{(q)}, \underline{h}) \right) \right\} \end{aligned}$$

$L(\mathcal{D}; \theta)$ is no longer convex \rightarrow convergence to a local minimum.

NLL gradient

$$\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = \sum_{q=1}^N \left\{ \sum_{y, \underline{h}} F_k(\underline{\mathbf{x}}^{(q)}, y, \underline{h}) p(y, \underline{h} | \underline{\mathbf{x}}; \theta) - \sum_{\underline{h}} F_k(\underline{\mathbf{x}}^{(q)}, y^{(q)}, \underline{h}) p(\underline{h} | y^{(q)}, \underline{\mathbf{x}}^{(q)}; \theta) \right\}$$

which can be again computed using the **forward-backward** method.

A gradient descent method (L-BFGS) can be again used to solve for θ .


NLL gradient

$$\frac{\partial L(\mathcal{D}; \theta)}{\partial \theta_k} = \sum_{q=1}^N \left\{ \sum_{y, \underline{h}} F_k(\underline{\mathbf{x}}^{(q)}, y, \underline{h}) p(y, \underline{h} | \underline{\mathbf{x}}; \theta) - \sum_{\underline{h}} F_k(\underline{\mathbf{x}}^{(q)}, y^{(q)}, \underline{h}) p(\underline{h} | y^{(q)}, \underline{\mathbf{x}}^{(q)}; \theta) \right\}$$

which can be again computed using the **forward-backward** method.

A gradient descent method (L-BFGS) can be again used to solve for θ .

Application to musical instrument classification



Feature functions for HCRF

Following (Quattoni et al., 2007)

$$\Psi(\underline{\mathbf{x}}, y, \underline{h}, \boldsymbol{\theta}) = \sum_{i=1}^N \langle \boldsymbol{\theta}(h_i), \mathbf{x}_i \rangle + \sum_{i=1}^N \theta(y, h_i) + \sum_{i=1}^N \theta(y, h_{i-1}, h_i)$$

- $\langle \boldsymbol{\theta}(h_i), \mathbf{x}_i \rangle$: compatibility between observation \mathbf{x}_i and hidden state $h_i \in \mathcal{H}$;
- $\theta(y, h_i)$: compatibility between hidden state h_i and label y ;
- $\theta(y, h_{i-1}, h_i)$: compatibility between transition $h_{i-1} \leftrightarrow h_i$ and label y .

Evaluation

- Classifying 1-second long segments of solo excerpts of **Cello, Guitar, Piano, Bassoon** and **Oboe**.
- Data:
 - training set: 2505 segments (*i.e.* 42');
 - testing set: 2505 segments.
- Classifiers:
 - ℓ_2 -regularized **HCRF** with 3 hidden states;
 - **Linear SVM**.
- Features: 47 cepstral, perceptual and temporal features.
- Results

Classifier	SVM	HCRF
Average accuracy	75%	76%

Evaluation

- Classifying 1-second long segments of solo excerpts of **Cello, Guitar, Piano, Bassoon** and **Oboe**.
- Data:
 - training set: 2505 segments (*i.e.* 42');
 - testing set: 2505 segments.
- Classifiers:
 - ℓ_2 -regularized **HCRF** with 3 hidden states;
 - **Linear SVM**.
- Features: 47 cepstral, perceptual and temporal features.
- Results


Classifier	SVM	HCRF
Average accuracy	75%	76%

- ▶ Introduction
- ▶ The logistic regression model
- ▶ Conditional Random Fields (for sequential data)
- ▶ Improvements and extensions to original CRFs
 - Regularization
 - Hidden-state CRF model
 - Other extensions
- ▶ Conclusion
- ▶ References
- ▶ Appendix

Other extensions

- **LDCRF**: Latent-Dynamic Conditional Random Field (Sung and Jurafsky, 2009)

modeling both hidden-states and structured-outputs.



Other extensions

- **LDCRF**: Latent-Dynamic Conditional Random Field (Morency et al., 2007)
modeling both hidden-states and structured-outputs.
- **Kernel-CRF** (Laferty et al., 2004; Altun et al., 2004)
introducing implicit features to account for (non-linear) interactions between original features.
- **Semi-Markov CRF** (Sarawagi and Cohen, 2005)
modeling segment-level labels
- **CCRF**: Continuous CRF (Qin and Liu, 2008)
modeling continuous labels in a regression setting.

Take-home messages

- CRFs are powerful structured-data prediction models (more flexible than HMMs and other more general Bayesian networks) as they are:
 - **discriminative models**: focus on modeling the target labels;
 - can handle a high number of feature functions, including transition features, and account for long-range dependencies.
 - **undirected models**: no need to normalize potentials locally.
 - allow for incorporating prior knowledge about constraints and label dependencies in an intuitive way.
- Easily **extendable** with key mechanisms: regularization, sparsity, latent variables, kernels...
- Great potential for various classification tasks (both symbolic and numerically-continuous data).

CRF software packages

Package	Language	Main features	Reference
CRF++	C++	Linear-chain CRF, NLP, L-BFGS optimization	(Taku-ku, 2003)
crfChain	Matlab, C mex	Linear-chain CRF, categorical features, L-BFGS optimization	Schmidt (2008)
CRFsuite	C++, Python	Linear-chain CRF, NLP, various regularization and optimization methods (L-BFGS), designed for fast training	(Okazaki, 2007)
HCRF library	C++, Matlab, Python	CRF, HCRF, LDCRF, continuous inputs, L-BFGS optimization	(Morency, 2010)
Mallet	Java	CRF, maxent, HMM, NLP, text feature extraction routines, various optimization methods (L-BFGS)	(McCallum, 2002)
Wapiti	C99	Linear-chain CRF, NLP, large label and feature sets, various regularization and optimization methods (L-BFGS, SGD), multi-threaded	(Lavergne et al., 2010)

CRF tutorials

- Charles Sutton and Andrew McCallum. **An Introduction to Conditional Random Fields for Relational Learning**. In *Introduction to Statistical Relational Learning*. Edited by Lise Getoor and Ben Taskar. MIT Press, 2006.
- Charles Elkan. **Log-linear Models and Conditional Random Fields**. *Notes for a tutorial at CIKM'08*, October 2008.
- Roman Klinger and Katrin Tomanek. **Classical Probabilistic Models and Conditional Random Fields**. *Algorithm Engineering Report TR07-2-2013*, December 2007.
- Hanna M. Wallach. **Conditional Random Fields: An Introduction**. Technical Report MS-CIS-04-21, Department of Computer and Information Science, University of Pennsylvania, 2004.
- Roland Memisevic. **An Introduction to Structured Discriminative Learning**. Technical Report, University of Toronto, 2006.
- Rahul Gupta. **Conditional Random Fields**. Unpublished report, IIT Bombay, 2006.

Bibliography I

- Y. Altun, A. Smola, and T. Hofmann. Exponential families for conditional random fields. In *Proceedings of the 20th conference on Uncertainty in artificial intelligence*, volume 0, 2004. URL <http://dl.acm.org/citation.cfm?id=1036844>.
- A. L. Berger, S. A. D. Pietra, and Vincent J. Della Pietra. A Maximum Entropy Approach to Natural Language Processing. *Computational Linguistics*, 22(1), 1996.
- L. Bottou. Stochastic Gradient Descent, 2010. URL <http://leon.bottou.org/projects/sgd>.
- M. Dunham and K. P. Murphy. pmtk3: probabilistic modeling toolkit for Matlab/Octave, version 3, 2010. URL <https://code.google.com/p/pmtk3/>.
- A. Gunawardana, M. Mahajan, A. Acero, and J. Platt. Hidden conditional random fields for phone classification. In *Interspeech*, 2005. URL <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.113.8190&rep=rep1&type=pdf>.
- T. Hastie, R. Tibshirani, and J. H. Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer, 2009. ISBN 0387848576. URL <http://books.google.com/books?id=tVIjMNS30b8C&pgis=1>.
- X. He, R. Zemel, and M. Carreira-Perpinan. Multiscale conditional random fields for image labeling. In *IEEE Computer Society Conference on Computer Vision and Pattern Recognition, CVPR*, volume 2, pages 695–702. IEEE, 2004. ISBN 0-7695-2158-4. doi: 10.1109/CVPR.2004.1315232. URL http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=1315232.
- W.-T. Hong. Speaker identification using Hidden Conditional Random Field-based speaker models. In *2010 International Conference on Machine Learning and Cybernetics*, pages 2811–2816. IEEE, July 2010. ISBN 978-1-4244-6526-2. doi: 10.1109/ICMLC.2010.5580793. URL http://ieeexplore.ieee.org/xpl/articleDetails.jsp?tp=&arnumber=5580793&contentType=Conference+Publications&searchField=Search_All&queryText=Hidden+Conditional+Random+Fields.
- V. Imbrasaite, T. Baltrusaitis, and P. Robinson. Emotion tracking in music using continuous conditional random fields and relative feature representation. In *2013 IEEE International Conference on Multimedia and Expo Workshops (ICMEW)*. IEEE, July 2013. ISBN 978-1-4799-1604-7. doi: 10.1109/ICMEW.2013.6618357. URL <http://ieeexplore.ieee.org/articleDetails.jsp?arnumber=6618357>.

Bibliography II

- F. V. Jensen and T. D. Nielsen. *Bayesian Networks and Decision Graphs*. Springer, 2007. ISBN 0387682813. URL http://books.google.fr/books/about/Bayesian_Networks_and_Decision_Graphs.html?id=goSLmQq4UBcC&pgis=1.
- C. Joder, S. Essid, and G. Richard. A Conditional Random Field Framework for Robust and Scalable Audio-to-Score Matching. *IEEE Transactions on Audio, Speech and Language Processing*, 19(8):2385–2397, 2011.
- J. Lafferty, Y. Liu, and X. Zhu. Kernel Conditional Random Fields: Representation, Clique Selection, and Semi-Supervised Learning. In *International Conference on Machine Learning*, 2004. URL <http://scholar.google.com/scholar?hl=en&btnG=Search&q=intitle:Kernel+Conditional+Random+Fields:+Representation+,+Clique+Selection+,+and+Semi-Supervised+Learning#5>.
- J. Lafferty, A. McCallum, and F. Pereira. Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *Machine learning international workshop then conference*, pages 282–289. Citeseer, 2001. URL <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.23.9849&rep=rep1&type=pdf>.
- T. Lavergne, O. Cappe, and F. Yvon. Practical very large scale CRFs. In *48th Annual Meeting of the Association for Computational Linguistics (ACL)*, number July, pages 504–513, Uppsala, Sweeden, 2010. URL <http://www.aclweb.org/anthology/P10-1052>.
- C.-X. Li. *Exploiting label correlations for multi-label classification*. Master thesis, UC San Diego, 2011.
- A. K. McCallum. Mallet: A machine learning for language toolkit, 2002. URL <http://mallet.cs.umass.edu>.
- L.-P. Morency. HCRF: Hidden- state Conditional Random Field Library, 2010. URL <http://sourceforge.net/projects/hcrf/>.
- L.-P. Morency, A. Quattoni, and T. Darrell. Latent-Dynamic Discriminative Models for Continuous Gesture Recognition. In *2007 IEEE Conference on Computer Vision and Pattern Recognition*, pages 1–8. IEEE, June 2007. ISBN 1-4244-1179-3. doi: 10.1109/CVPR.2007.383299. URL <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=4270324>.
- J. Morris and E. Fosler-Lussier. Conditional Random Fields for Integrating Local Discriminative Classifiers. *TASLP*, 16(3):617–628, 2008. ISSN 1558-7916. doi: 10.1109/TASL.2008.916057. URL <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=4451149>.

Bibliography III

- N. Okazaki. CRFsuite: a fast implementation of Conditional Random Fields (CRFs), 2007. URL <http://www.chokkan.org/software/crfsuite/>.
- S. Ortmanns, H. Ney, and A. Eiden. Language-model look-ahead for large vocabulary speech recognition. In *Proceeding of Fourth International Conference on Spoken Language Processing. ICSLP '96*, volume 4, pages 2095–2098. IEEE, 1996. ISBN 0-7803-3555-4. doi: 10.1109/ICSLP.1996.607215. URL <http://ieeexplore.ieee.org/articleDetails.jsp?arnumber=607215>.
- T. Qin and T. Liu. Global ranking using continuous conditional random fields. In *Advances in Neural and Information Processing Systems*, 2008. URL http://machinelearning.wustl.edu/mlpapers/paper_files/NIPS2008_0518.pdf.
- A. Quattoni, S. Wang, L.-P. Morency, M. Collins, and T. Darrell. Hidden conditional random fields. *IEEE transactions on pattern analysis and machine intelligence*, 29(10):1848–53, Oct. 2007. ISSN 0162-8828. doi: 10.1109/TPAMI.2007.1124. URL <http://www.ncbi.nlm.nih.gov/pubmed/17699927>.
- S. Reiter, B. Schuller, and G. Rigoll. Hidden Conditional Random Fields for Meeting Segmentation. In *Multimedia and Expo, 2007 IEEE International Conference on*, pages 639–642. IEEE, July 2007. ISBN 1-4244-1016-9. doi: 10.1109/ICME.2007.4284731. URL http://ieeexplore.ieee.org/xpl/articleDetails.jsp?tp=&arnumber=4284731&contentType=Conference+Publications&searchField=Search_All&queryText=Hidden+Conditional+Random+Fields.
- O. Rudovic, V. Pavlovic, and M. Pantic. Kernel conditional ordinal random fields for temporal segmentation of facial action units. In *Computer Vision - ECCV*, pages 1–10, 2012. URL http://link.springer.com/chapter/10.1007/978-3-642-33868-7_26.
- S. Sarawagi and W. Cohen. Semi-markov conditional random fields for information extraction. In *Advances in Neural Information Processing Systems*, 2005. URL http://machinelearning.wustl.edu/mlpapers/paper_files/NIPS2005_427.pdf.
- E. Schmidt and Y. Kim. Modeling musical emotion dynamics with conditional random fields. In *International Conference on Music Information Retrieval (ISMIR)*, number Ismir, pages 777–782, Miami, FL, USA, 2011. URL <http://music.ece.drexel.edu/files/Navigation/Publications/Schmidt2011c.pdf>.
- M. Schmidt. crfChain, 2008. URL <http://www.di.ens.fr/~mschmidt/Software/crfChain.html>.

Bibliography IV

- B. Settles. Biomedical named entity recognition using conditional random fields and rich feature sets. *International Joint Workshop on Natural Language Processing in Biomedicine and its Applications*, 2004. URL <http://dl.acm.org/citation.cfm?id=1567618>.
- N. Sokolovska. *Contributions to the estimation of probabilistic discriminative models: semi-supervised learning and feature selection*. PhD thesis, Telecom ParisTech, 2010.
- Y.-H. Sung and D. Jurafsky. Hidden Conditional Random Fields for phone recognition. *2009 IEEE Workshop on Automatic Speech Recognition & Understanding*, pages 107–112, Dec. 2009. doi: 10.1109/ASRU.2009.5373329. URL <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=5373329>.
- Taku-ku. CRF++, 2003. URL <http://crfpp.googlecode.com/svn/trunk/doc/index.html>.
- B. Taskar, P. Abbeel, and D. Koller. Discriminative probabilistic models for relational data. In *Proceedings of the Eighteenth conference on Uncertainty in Artificial Intelligence (UAI02)*, 2002. URL <http://dl.acm.org/citation.cfm?id=2073934>.
- H. M. Wallach. *Efficient Training of Conditional Random Fields*. PhD thesis, University of Edinburgh, 2002.
- S. B. Wang, A. Quattoni, L.-P. Morency, and D. Demirdjian. Hidden Conditional Random Fields for Gesture Recognition. pages 1521–1527. IEEE Computer Society, 2006. ISBN 0-7695-2597-0. URL <http://portal.acm.org/citation.cfm?id=1153171.1153627>.
- G. Yuan, K. Chang, C. Hsieh, and C. Lin. A comparison of optimization methods and software for large-scale l1-regularized linear classification. *The Journal of Machine Learning Research*, 2010. URL <http://dl.acm.org/citation.cfm?id=1953034>.
- J. Zhu and T. Hastie. Support vector machines, kernel logistic regression and boosting. *Multiple Classifier Systems*, 2364:16–26, 2002. URL http://link.springer.com/chapter/10.1007/3-540-45428-4_2.

► Improvements and extensions to original CRFs

► Conclusion

► References

► **Appendix**

- Optimization with stochastic gradient learning
- Comparing KLR and SVM
- Derivation of the Viterbi algorithm
- The forward-backward method

LR model learning with stochastic gradient descent (SGL)

- **Idea:** make gradient updates based on **one** training example **at a time**
- **Use:** $\frac{\partial L(\mathcal{D}; \tilde{\mathbf{w}})}{\partial w_j} = (y_i - p(\mathbf{x}_i; \tilde{\mathbf{w}})) x_{ji}$

Algorithm

```

- Initialise  $\tilde{\mathbf{w}}$ 
- Repeat (until convergence)
    - Randomly permute training examples  $\mathbf{x}_i$ 
    - For  $i = 1 : N$ 
         $w_j \leftarrow w_j + t (y_{\sigma_i} - p_{\sigma_i}) x_{j\sigma_i} ; j = 1, \dots, D$ 

```

- t : *step size*, to be tuned
- **Complexity of SGL:** $O(NFD)$ per *epoch*; with F the average number of non-zero feature coefficients per example; an *epoch* is a “complete” update using all training examples.

LR model learning with stochastic gradient descent (SGL)

- **Idea:** make gradient updates based on **one** training example **at a time**
- **Use:** $\frac{\partial L(\mathcal{D}; \tilde{\mathbf{w}})}{\partial w_j} = (y_i - p(\mathbf{x}_i; \tilde{\mathbf{w}})) x_{ji}$

Algorithm

```

- Initialise  $\tilde{\mathbf{w}}$ 
- Repeat (until convergence)
    - Randomly permute training examples  $\mathbf{x}_i$ 
    - For  $i = 1 : N$ 
         $w_j \leftarrow w_j + t (y_{\sigma_i} - p_{\sigma_i}) x_{j\sigma_i} ; j = 1, \dots, D$ 

```

- t : *step size*, to be tuned
- **Complexity of SGL:** $O(NFD)$ per *epoch*; with F the average number of non-zero feature coefficients per example; an *epoch* is a “complete” update using all training examples.

LR model learning with stochastic gradient descent (SGL)

- **Idea:** make gradient updates based on **one** training example **at a time**
- **Use:** $\frac{\partial L(\mathcal{D}; \tilde{\mathbf{w}})}{\partial w_j} = (y_i - p(\mathbf{x}_i; \tilde{\mathbf{w}})) x_{ji}$

Algorithm

```

- Initialise  $\tilde{\mathbf{w}}$ 
- Repeat (until convergence)
    - Randomly permute training examples  $\mathbf{x}_i$ 
    - For  $i = 1 : N$ 
         $w_j \leftarrow w_j + t (y_{\sigma_i} - p_{\sigma_i}) x_{j\sigma_i} ; j = 1, \dots, D$ 

```

- t : *step size*, to be tuned
- **Complexity of SGL:** $O(NFD)$ per *epoch*; with F the average number of non-zero feature coefficients per example; an *epoch* is a “complete” update using all training examples.

Support Vector Machines


Recalling SVM as a regularized function fitting problem

- The SVM solution, $g(\mathbf{x}) = w_0 + \mathbf{w}^T \phi(\mathbf{x})$, can be found by solving:

$$\min_{\tilde{\mathbf{w}}} \sum_{i=1}^N [1 - y'_i g(\mathbf{x}_i)]_+ + \frac{\gamma}{2} \|g\|_{\mathcal{H}_K}^2 ; y'_i \in \{-1, 1\}$$

Hinge loss

$$[1 - y'_i g(\mathbf{x}_i)]_+ = \max(0, 1 - y'_i g(\mathbf{x}_i))$$



KLR vs SVM

- Let $y'_i = \begin{cases} 1 & \text{if } y_i = 1 \\ -1 & \text{if } y_i = 0 \end{cases}$
- The negative log-likelihood of the KLR model can then be written as $L(\mathcal{D}; \tilde{\mathbf{w}}) = \sum_{i=1}^N \log(1 + \exp -y'_i g(\mathbf{x}_i))$.
- Both KLR and SVM solve:

$$\min_{\tilde{\mathbf{w}}} \sum_{i=1}^N l(y'_i g(\mathbf{x}_i)) + \frac{\lambda}{2} \|\mathbf{g}\|_{\mathcal{H}_{\mathcal{K}}}^2;$$

KLR


$$l(y'_i g(\mathbf{x}_i)) = \log(1 + \exp -y'_i f(\mathbf{x}_i))$$

SVM

$$l(y'_i g(\mathbf{x}_i)) = [1 - y'_i f(\mathbf{x}_i)]_+$$

KLR vs SVM

Hinge vs negative binomial log-likelihood



Decoding the optimal sequence

$$\text{Let: } g_i(y_{i-1}, y_i) \triangleq \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i); \text{ then:}$$

$$\hat{y} = \underset{\underline{y}}{\operatorname{argmax}} \sum_{i=1}^n \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i) = \underset{\underline{y}}{\operatorname{argmax}} \sum_{i=1}^n g_i(y_{i-1}, y_i).$$

$$\begin{aligned} \text{Let } \delta_m(\underline{s}) &\triangleq \max_{\{y_1, \dots, y_{m-1}\}} \left[\sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + g_m(y_{m-1}, \underline{s}) \right] \\ &= \max_{\{y_1, \dots, y_{m-1}\}} \sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + \max_{y_{m-1}} g_m(y_{m-1}, \underline{s}). \end{aligned}$$

Decoding the optimal sequence

$$\text{Let: } g_i(y_{i-1}, y_i) \triangleq \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i); \text{ then:}$$

$$\hat{y} = \underset{\underline{y}}{\operatorname{argmax}} \sum_{i=1}^n \sum_{j=1}^D \theta_j f_j(y_{i-1}, y_i, \underline{\mathbf{x}}, i) = \underset{\underline{y}}{\operatorname{argmax}} \sum_{i=1}^n g_i(y_{i-1}, y_i).$$

$$\begin{aligned} \text{Let } \delta_m(\underline{s}) &\triangleq \max_{\{y_1, \dots, y_{m-1}\}} \left[\sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + g_m(y_{m-1}, \underline{s}) \right] \\ &= \max_{\{y_1, \dots, y_{m-1}\}} \sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + \max_{y_{m-1}} g_m(y_{m-1}, \underline{s}). \end{aligned}$$

Viterbi decoding

$$\delta_m(s) = \max_{\{y_1, \dots, y_{m-1}\}} \sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + \max_{y_{m-1}} g_m(y_{m-1}, s)$$

$$\begin{aligned} \text{So: } \delta_{m-1}(y_{m-1}) &= \max_{\{y_1, \dots, y_{m-2}\}} \left[\sum_{i=1}^{m-2} g_i(y_{i-1}, y_i) + g_{m-1}(y_{m-2}, y_{m-1}) \right] \\ &= \max_{\{y_1, \dots, y_{m-2}\}} \sum_{i=1}^{m-1} g_i(y_{i-1}, y_i). \end{aligned}$$

$$\delta_m(s) = \max_{y_{m-1} \in \mathcal{Y}} [\delta_{m-1}(y_{m-1}) + g_m(y_{m-1}, s)].$$

Viterbi decoding

$$\delta_m(s) = \max_{\{y_1, \dots, y_{m-1}\}} \sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + \max_{y_{m-1}} g_m(y_{m-1}, s)$$

$$\begin{aligned} \text{So: } \delta_{m-1}(y_{m-1}) &= \max_{\{y_1, \dots, y_{m-2}\}} \left[\sum_{i=1}^{m-2} g_i(y_{i-1}, y_i) + g_{m-1}(y_{m-2}, y_{m-1}) \right] \\ &= \max_{\{y_1, \dots, y_{m-2}\}} \sum_{i=1}^{m-1} g_i(y_{i-1}, y_i). \end{aligned}$$

$$\delta_m(s) = \max_{y_{m-1} \in \mathcal{Y}} [\delta_{m-1}(y_{m-1}) + g_m(y_{m-1}, s)].$$

Viterbi decoding

The algorithm

Initialization:

$$\delta_1(s) = g_1(y_0, s); \forall s \in \mathcal{Y}; y_0 = \text{start}$$

$$\psi_1(s) = \text{start}$$

Recursion:

$$\forall s \in \mathcal{Y}; 1 \leq m \leq n$$

$$\delta_m(s) = \max_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

$$\psi_m(s) = \operatorname{argmax}_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

Termination:

$$\delta_n(y_n^*) = \max_{y \in \mathcal{Y}} \delta_n(y)$$

$$y_n^* = \operatorname{argmax}_{y \in \mathcal{Y}} \delta_n(y)$$

Path backtracking:

$$y_m^* = \psi_{m+1}(y_{m+1}^*); m = n-1, n-2, \dots, 1.$$

Viterbi decoding

The algorithm

Initialization:

$$\delta_1(s) = g_1(y_0, s); \forall s \in \mathcal{Y}; y_0 = \text{start}$$

$$\psi_1(s) = \text{start}$$

Recursion:

$$\forall s \in \mathcal{Y}; 1 \leq m \leq n$$

$$\delta_m(s) = \max_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

$$\psi_m(s) = \operatorname{argmax}_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

Termination:

$$\delta_n(y_n^*) = \max_{y \in \mathcal{Y}} \delta_n(y)$$

$$y_n^* = \operatorname{argmax}_{y \in \mathcal{Y}} \delta_n(y)$$

Path backtracking:

$$y_m^* = \psi_{m+1}(y_{m+1}^*); m = n-1, n-2, \dots, 1.$$

Viterbi decoding

The algorithm

Initialization:

$$\delta_1(s) = g_1(y_0, s); \forall s \in \mathcal{Y}; y_0 = \text{start}$$

$$\psi_1(s) = \text{start}$$

Recursion:

$$\forall s \in \mathcal{Y}; 1 \leq m \leq n$$

$$\delta_m(s) = \max_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

$$\psi_m(s) = \operatorname{argmax}_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

Termination:

$$\delta_n(y_n^*) = \max_{y \in \mathcal{Y}} \delta_n(y)$$

$$y_n^* = \operatorname{argmax}_{y \in \mathcal{Y}} \delta_n(y)$$

Path backtracking:

$$y_m^* = \psi_{m+1}(y_{m+1}^*); m = n-1, n-2, \dots, 1.$$

Viterbi decoding

The algorithm

Initialization:

$$\delta_1(s) = g_1(y_0, s); \forall s \in \mathcal{Y}; y_0 = \text{start}$$

$$\psi_1(s) = \text{start}$$

Recursion:

$$\forall s \in \mathcal{Y}; 1 \leq m \leq n$$

$$\delta_m(s) = \max_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

$$\psi_m(s) = \operatorname{argmax}_{y \in \mathcal{Y}} [\delta_{m-1}(y) + g_m(y, s)]$$

Termination:

$$\delta_n(y_n^*) = \max_{y \in \mathcal{Y}} \delta_n(y)$$

$$y_n^* = \operatorname{argmax}_{y \in \mathcal{Y}} \delta_n(y)$$

Path backtracking:

$$y_m^* = \psi_{m+1}(y_{m+1}^*); m = n-1, n-2, \dots, 1.$$

The forward recursion

Define α scores as:

$$\alpha_1(y_1) = M_1(y_0, y_1)$$

$$\alpha_2(y_2) = \sum_{y_1 \in \mathcal{Y}} M_2(y_1, y_2) \alpha_1(y_1)$$

$$\alpha_3(y_3) = \sum_{y_2 \in \mathcal{Y}} M_3(y_2, y_3) \alpha_2(y_2)$$

$$\vdots$$

$$\alpha_m(y_m) = \sum_{y_{m-1}} M_m(y_{m-1}, y_m) \alpha_{m-1}(y_{m-1}); \quad 2 \leq m \leq n$$

The forward recursion

Define α scores as:

$$\alpha_1(y_1) = M_1(y_0, y_1)$$

$$\alpha_2(y_2) = \sum_{y_1 \in \mathcal{Y}} M_2(y_1, y_2) \alpha_1(y_1)$$

$$\alpha_3(y_3) = \sum_{y_2 \in \mathcal{Y}} M_3(y_2, y_3) \alpha_2(y_2) = \sum_{y_1, y_2} M_3(y_2, y_3) M_2(y_1, y_2) M_1(y_0, y_1)$$

\vdots

$$\alpha_m(y_m) = \sum_{y_{m-1}} M_m(y_{m-1}, y_m) \alpha_{m-1}(y_{m-1}); \quad 2 \leq m \leq n$$

At the end of the sequence

$$\sum_{y_n \in \mathcal{Y}} \alpha_n(y_n) = \sum_{\underline{y} \in \mathcal{Y}^n} \prod_{i=1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) = Z(\underline{\mathbf{x}}, \boldsymbol{\theta}).$$

The forward recursion

Define α scores as:

$$\alpha_m(y_m) = \sum_{y_{m-1}} M_m(y_{m-1}, y_m) \alpha_{m-1}(y_{m-1}); \quad 2 \leq m \leq n.$$

At the end of the sequence

$$\sum_{y_n \in \mathcal{Y}} \alpha_n(y_n) = \sum_{\underline{y} \in \mathcal{Y}^n} \prod_{i=1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) = Z(\underline{\mathbf{x}}, \boldsymbol{\theta}).$$

Complexity: reduced from $O(K^n)$ to $O(nK^2)$.

The backward recursion

$$\begin{aligned}\beta_m(y_m) &= \sum_{y_{m+1} \in \mathcal{Y}} M_{m+1}(y_m, y_{m+1}) \beta_{m+1}(y_{m+1}); \quad 1 \leq m \leq n-1 \\ \beta_n(y_n) &= 1\end{aligned}$$

At the beginning of the sequence

$$Z(\underline{\mathbf{x}}, \boldsymbol{\theta}) = \sum_{y_1 \in \mathcal{Y}} M_1(y_0, y_1) \beta_1(y_1).$$

Marginal probability

$$p(y_{m-1}, y_m | \underline{\mathbf{x}}) = \sum_{\underline{y} \setminus \{y_{m-1}, y_m\}} p(\underline{y} | \underline{\mathbf{x}});$$

$$\underline{y} \setminus \{y_{m-1}, y_m\} \triangleq \{y_1, \dots, y_{\textcolor{red}{m}-2}, y_{\textcolor{red}{m}+1}, \dots, y_n\}.$$

$$\begin{aligned} p(y_{m-1}, y_m | \underline{\mathbf{x}}) &= \frac{1}{Z(\underline{\mathbf{x}})} \sum_{\underline{y} \setminus \{y_{m-1}, y_m\}} \prod_{i=1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) \\ &= \frac{1}{Z(\underline{\mathbf{x}})} \sum_{\underline{y} \setminus \{y_{m-1}, y_m\}} \prod_{i=1}^{\textcolor{red}{m}-1} M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) \times M_{\textcolor{red}{m}}(y_{m-1}, y_m, \underline{\mathbf{x}}) \\ &\quad \times \prod_{i=\textcolor{red}{m}+1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) \end{aligned}$$

Marginal probability

$$p(y_{m-1}, y_m | \underline{\mathbf{x}}) = \sum_{\underline{y} \setminus \{y_{m-1}, y_m\}} p(\underline{y} | \underline{\mathbf{x}});$$

$$\underline{y} \setminus \{y_{m-1}, y_m\} \triangleq \{y_1, \dots, y_{\textcolor{red}{m}-2}, y_{\textcolor{red}{m}+1}, \dots, y_n\}.$$

$$\begin{aligned} p(y_{m-1}, y_m | \underline{\mathbf{x}}) &= \frac{1}{Z(\underline{\mathbf{x}})} \sum_{\underline{y} \setminus \{y_{m-1}, y_m\}} \prod_{i=1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) \\ &= \frac{1}{Z(\underline{\mathbf{x}})} \sum_{\underline{y} \setminus \{y_{m-1}, y_m\}} \prod_{i=1}^{\textcolor{red}{m}-1} M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) \times M_{\textcolor{red}{m}}(y_{m-1}, y_m, \underline{\mathbf{x}}) \\ &\quad \times \prod_{i=\textcolor{red}{m}+1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) \end{aligned}$$

Marginal probability

$$\begin{aligned}
 p(y_{m-1}, y_m | \underline{\mathbf{x}}) &= \frac{1}{Z(\underline{\mathbf{x}})} M_{\textcolor{red}{m}}(y_{m-1}, y_m, \underline{\mathbf{x}}) \times \sum_{\{y_1, \dots, y_{\textcolor{red}{m}-2}\}} \prod_{i=1}^{\textcolor{red}{m}-1} M_i(y_{i-1}, y_i, \underline{\mathbf{x}}) \\
 &\quad \times \sum_{\{y_{\textcolor{red}{m}+1}, \dots, y_n\}} \prod_{i=\textcolor{red}{m}+1}^n M_i(y_{i-1}, y_i, \underline{\mathbf{x}})
 \end{aligned}$$

$$p(y_{m-1}, y_m | \underline{\mathbf{x}}) = \frac{1}{Z(\underline{\mathbf{x}})} \alpha_{\textcolor{red}{m}-1}(y_{m-1}) M_{\textcolor{red}{m}}(y_{m-1}, y_m, \underline{\mathbf{x}}) \beta_{\textcolor{red}{m}}(y_m).$$